

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: Columns = ([ 'duration', 'protocol_type', 'service', 'flag', 'src_bytes', 'dst_bytes', 'land', 'wrong_fragment', 'urgent', 'hot',
    'num_failed_logins', 'logged_in', 'num_compromised', 'root_shell', 'su_attempted', 'num_root', 'num_file_creations',
    'num_shells', 'num_access_files', 'num_outbound_cmds', 'is_host_login', 'is_guest_login', 'count', 'srv_count',
    'serror_rate', 'srv_serror_rate', 'rerror_rate', 'srv_rerror_rate', 'same_srv_rate', 'diff_srv_rate', 'srv_diff_host_rate',
    'dst_host_count', 'dst_host_srv_count', 'dst_host_same_srv_rate', 'dst_host_diff_srv_rate', 'dst_host_same_src_port_rate',
    'dst_host_srv_diff_host_rate', 'dst_host_serror_rate', 'dst_host_srv_serror_rate', 'dst_host_rerror_rate',
    'dst_host_srv_rerror_rate', 'attack', 'level'])
```

```
In [4]: data = pd.read_csv("nsl-kdd/KDDTrain+.txt", sep = ",", encoding = 'utf-8')
```

```
In [5]: # Load data
data.columns = Columns
```

```
In [6]: null_counts = data.isnull().sum()
# Print the number of null values
print(f"{null_counts.sum()} null entries have been found in the dataset\n")
# Drop null values
data.dropna(inplace=True) # or df_data = df_data.dropna()
```

```
# Find and handle duplicates
duplicate_count = data.duplicated().sum()
# Print the number of duplicate entries
print(f"{duplicate_count} duplicate entries have been found in the dataset\n")
# Remove duplicates
data.drop_duplicates(inplace=True) # or df_data = df_data.drop_duplicates()
# Display relative message
print(f"All duplicates have been removed\n")
```

```
# Reset the indexes
data.reset_index(drop=True, inplace=True)

# Inspect the dataset for categorical columns
print("Categorical columns:", data.select_dtypes(include=['object']).columns.tolist(), '\n')

# Print the first 5 Lines
data.head()
```

0 null entries have been found in the dataset

0 duplicate entries have been found in the dataset

All duplicates have been removed

Categorical columns: ['protocol_type', 'service', 'flag', 'attack']

```
Out[6]: duration protocol_type service flag src_bytes dst_bytes land wrong_fragment urgent hot ... dst_host_same_srv_rate dst_host_diff_srv_rate dst_host_same_src_port_rate dst_host_srv_diff_host_rate dst_host_serr...
```

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	...	dst_host_same_srv_rate	dst_host_diff_srv_rate	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate	dst_host_serr...
0	0	udp	other	SF	146	0	0	0	0	0	...	0.00	0.60	0.88	0.00	0.00
1	0	tcp	private	S0	0	0	0	0	0	0	...	0.10	0.05	0.00	0.00	0.00
2	0	tcp	http	SF	232	8153	0	0	0	0	0	1.00	0.00	0.03	0.04	0.00
3	0	tcp	http	SF	199	420	0	0	0	0	0	1.00	0.00	0.00	0.00	0.00
4	0	tcp	private	REJ	0	0	0	0	0	0	0	0.07	0.07	0.00	0.00	0.00

5 rows × 43 columns

```
In [7]: # number of attack Labels
data['attack'].value_counts()
```

```
Out[7]: normal      67342
neptune     41214
satan       3633
ipsweep     3599
portsweep   2931
smurf       2646
nmap        1493
back        956
teardrop    892
warezclient 890
pod         201
guess_passwd 53
buffer_overflow 30
warezmaster  20
land        18
imap        11
rootkit     10
loadmodule   9
ftp_write    8
multihop     7
phf         4
perl         3
spy          2
Name: attack, dtype: int64
```

```
In [8]: # changing attack Labels to their respective attack class
def change_label(df):
    df.attack.replace(['back', 'land', 'neptune', 'pod', 'smurf', 'teardrop'], 'Dos', inplace=True)
    df.attack.replace(['guess_passwd', 'imap', 'ftp_write', 'multihop', 'phf', 'spy', 'warezclient', 'warezmaster'], 'R2L', inplace=True)
    df.attack.replace(['ipsweep', 'nmap', 'portsweep', 'satan'], 'Probe', inplace=True)
    df.attack.replace(['buffer_overflow', 'loadmodule', 'perl', 'rootkit'], 'U2R', inplace=True)
```

```
In [9]: change_label(data)
```

```
In [10]: # distribution of attack classes
data.attack.value_counts()
```

```
Out[10]: normal      67342
Dos        45927
Probe      11656
R2L        995
U2R        52
Name: attack, dtype: int64
```

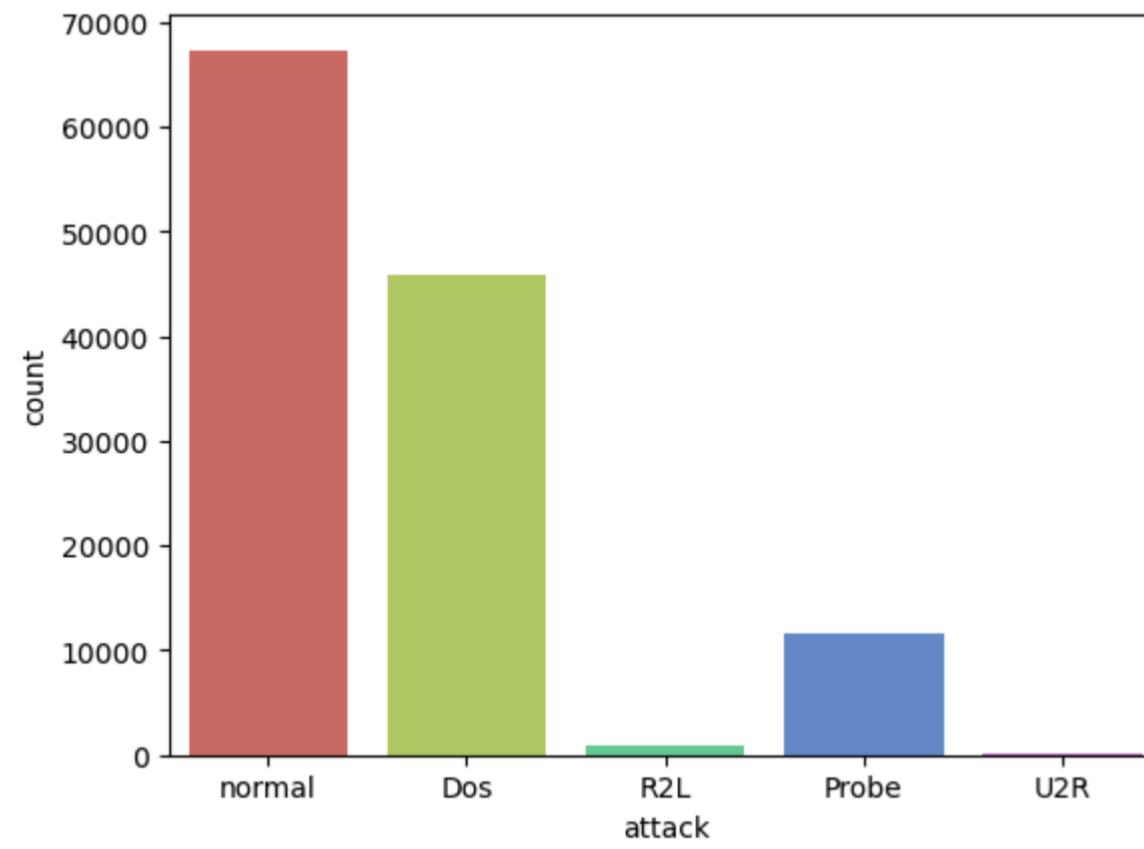
In [11]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 125972 entries, 0 to 125971
Data columns (total 43 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   duration         125972 non-null   int64  
 1   protocol_type    125972 non-null   object  
 2   service          125972 non-null   object  
 3   flag             125972 non-null   object  
 4   src_bytes        125972 non-null   int64  
 5   dst_bytes        125972 non-null   int64  
 6   land             125972 non-null   int64  
 7   wrong_fragment   125972 non-null   int64  
 8   urgent            125972 non-null   int64  
 9   hot               125972 non-null   int64  
 10  num_failed_logins 125972 non-null   int64  
 11  logged_in        125972 non-null   int64  
 12  num_compromised  125972 non-null   int64  
 13  root_shell       125972 non-null   int64  
 14  su_attempted     125972 non-null   int64  
 15  num_root          125972 non-null   int64  
 16  num_file_creations 125972 non-null   int64  
 17  num_shells        125972 non-null   int64  
 18  num_access_files  125972 non-null   int64  
 19  num_outbound_cmds 125972 non-null   int64  
 20  is_host_login    125972 non-null   int64  
 21  is_guest_login   125972 non-null   int64  
 22  count              125972 non-null   int64  
 23  srv_count         125972 non-null   int64  
 24  serror_rate       125972 non-null   float64 
 25  srv_serror_rate   125972 non-null   float64 
 26  rerror_rate       125972 non-null   float64 
 27  srv_error_rate   125972 non-null   float64 
 28  same_srv_rate    125972 non-null   float64 
 29  diff_srv_rate    125972 non-null   float64 
 30  srv_diff_host_rate 125972 non-null   float64 
 31  dst_host_count   125972 non-null   int64  
 32  dst_host_srv_count 125972 non-null   int64  
 33  dst_host_same_srv_rate 125972 non-null   float64 
 34  dst_host_diff_srv_rate 125972 non-null   float64 
 35  dst_host_same_src_port_rate 125972 non-null   float64 
 36  dst_host_srv_diff_host_rate 125972 non-null   float64 
 37  dst_host_error_rate 125972 non-null   float64 
 38  dst_host_srv_error_rate 125972 non-null   float64 
 39  dst_host_rerror_rate 125972 non-null   float64 
 40  dst_host_srv_rerror_rate 125972 non-null   float64 
 41  attack             125972 non-null   object  
 42  level              125972 non-null   int64 

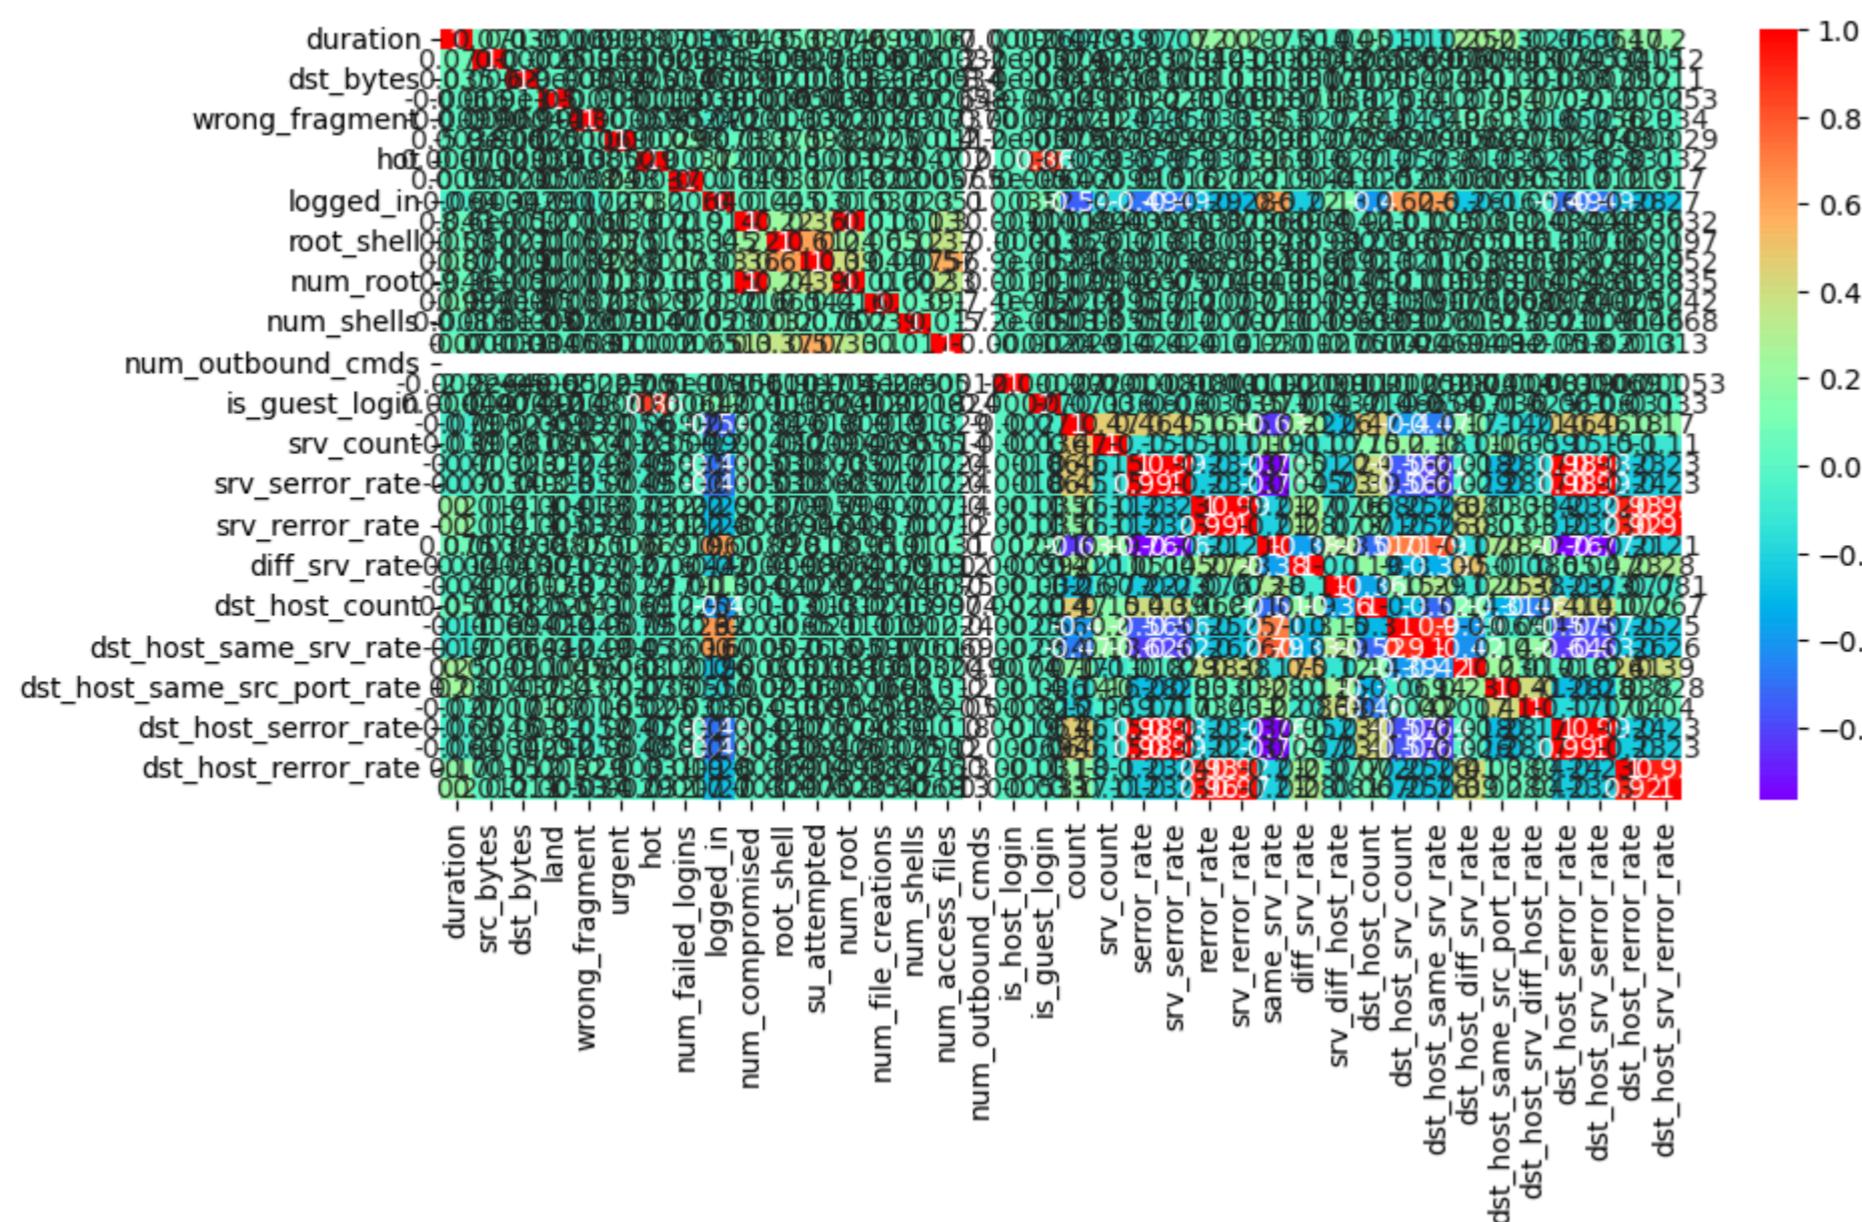
dtypes: float64(15), int64(24), object(4)
memory usage: 41.3+ MB
```

In [13]: `del data['level']`

```
In [12]: sns.countplot(x='attack', data=data, palette='hls')
plt.show()
#plt.savefig('count_plot') mal: the nodule malignancy, 0: benign, 1: malignant
```



```
In [14]: plt.figure(figsize = (10,5))
sns.heatmap(data.corr(), annot = True, cmap="rainbow")
plt.show()
```



```
In [16]: # Import Label encoder
from sklearn import preprocessing

# Label_encoder object knows
# how to understand word labels.
label_encoder = preprocessing.LabelEncoder()

# Encode Labels in column 'species'.
data['attack']= label_encoder.fit_transform(data['attack'])
data['protocol_type']= label_encoder.fit_transform(data['protocol_type'])
data['service']= label_encoder.fit_transform(data['service'])
data['flag']= label_encoder.fit_transform(data['flag'])
```

```
In [17]: X = data.drop(["attack"],axis =1)
y = data["attack"]
```

FS

```
In [18]: from sklearn.feature_selection import SelectKBest, SelectPercentile, mutual_info_classif
```

```
In [19]: selector = SelectPercentile(mutual_info_classif, percentile=25)
X_reduced = selector.fit_transform(X, y)
#X_reduced.shape
```

```
In [20]: cols = selector.get_support(indices=True)
selected_columns = X.iloc[:,cols].columns.tolist()
selected_columns
```

```
Out[20]: ['service',
'flag',
'src_bytes',
'dst_bytes',
'logged_in',
'count',
'same_srv_rate',
'diff_srv_rate',
'dst_host_srv_count',
'dst_host_diff_srv_rate']
```

```
In [21]: len(selected_columns)
```

```
Out[21]: 10
```

```
In [22]: df = data[['service',
'flag',
'src_bytes',
'dst_bytes',
'logged_in',
'count',
'same_srv_rate',
'diff_srv_rate',
'dst_host_srv_count',
'dst_host_diff_srv_rate','attack']]
```

```
In [23]: df.to_csv('ns1_processed.csv')
```

```
In [3]: df = pd.read_csv('ns1_processed.csv')
```

```
In [4]: del df['Unnamed: 0']
```

```
In [5]: df.columns
```

```
Out[5]: Index(['service', 'flag', 'src_bytes', 'dst_bytes', 'logged_in', 'count',
'same_srv_rate', 'diff_srv_rate', 'dst_host_srv_count',
'dst_host_diff_srv_rate', 'attack'],
dtype='object')
```

```
In [6]: X = df[['service', 'flag', 'src_bytes', 'dst_bytes', 'logged_in', 'count',
'same_srv_rate', 'diff_srv_rate', 'dst_host_srv_count',
'dst_host_diff_srv_rate',]]
y = df["attack"]
```

```
In [7]: from sklearn.model_selection import train_test_split
```

```
In [8]: from sklearn.metrics import accuracy_score # for calculating accuracy of model
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
In [9]: ML_Model = []
accuracy = []
precision = []
recall = []

f1score = []

#function to call for storing the results
def storeResults(model, a,b,c,d):
    ML_Model.append(model)
    accuracy.append(round(a, 3))
    precision.append(round(b, 3))
    recall.append(round(c, 3))
    f1score.append(round(d, 3))
```

```
In [10]: # importing lime
import lime
from lime import lime_tabular
import shap
```

DNN

```
In [11]: from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)
# Scale the features using StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [13]: # Build the model architecture
model = tf.keras.Sequential([
    tf.keras.layers.Dense(1024, activation='relu', input_dim=X_train.shape[1]),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(256, activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(128, activation='tanh'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])

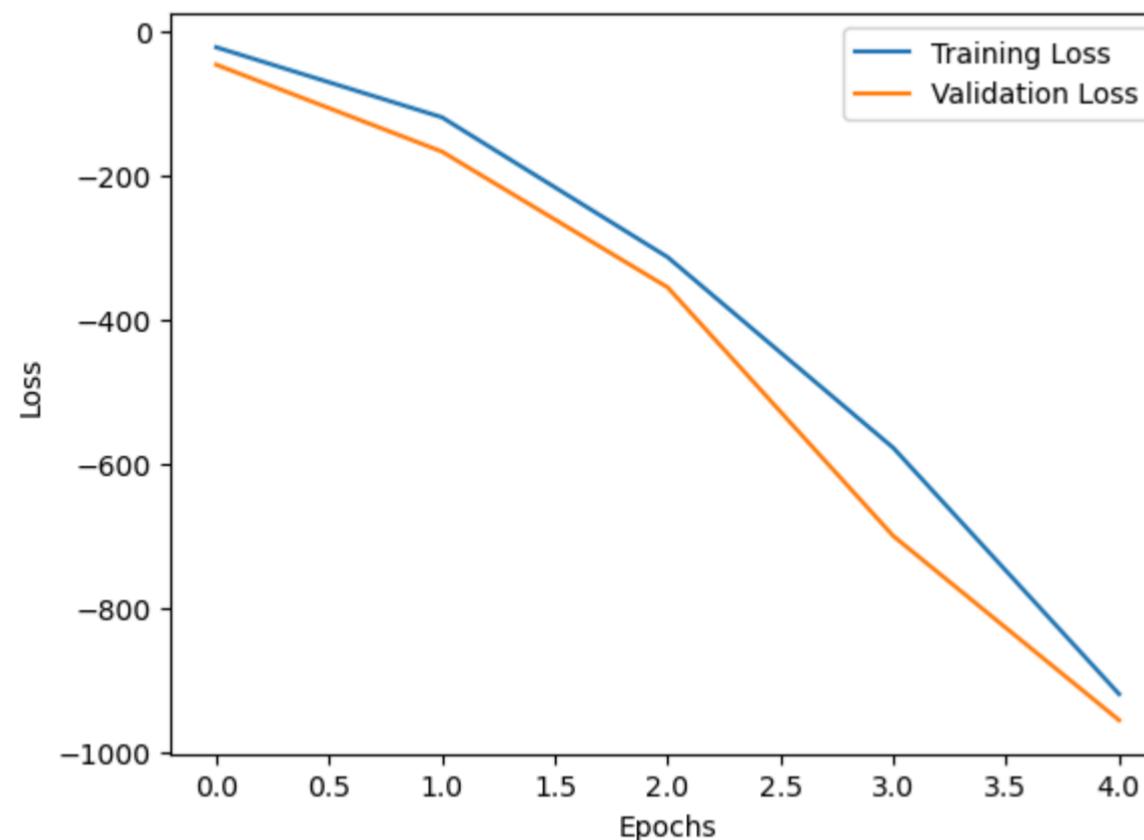
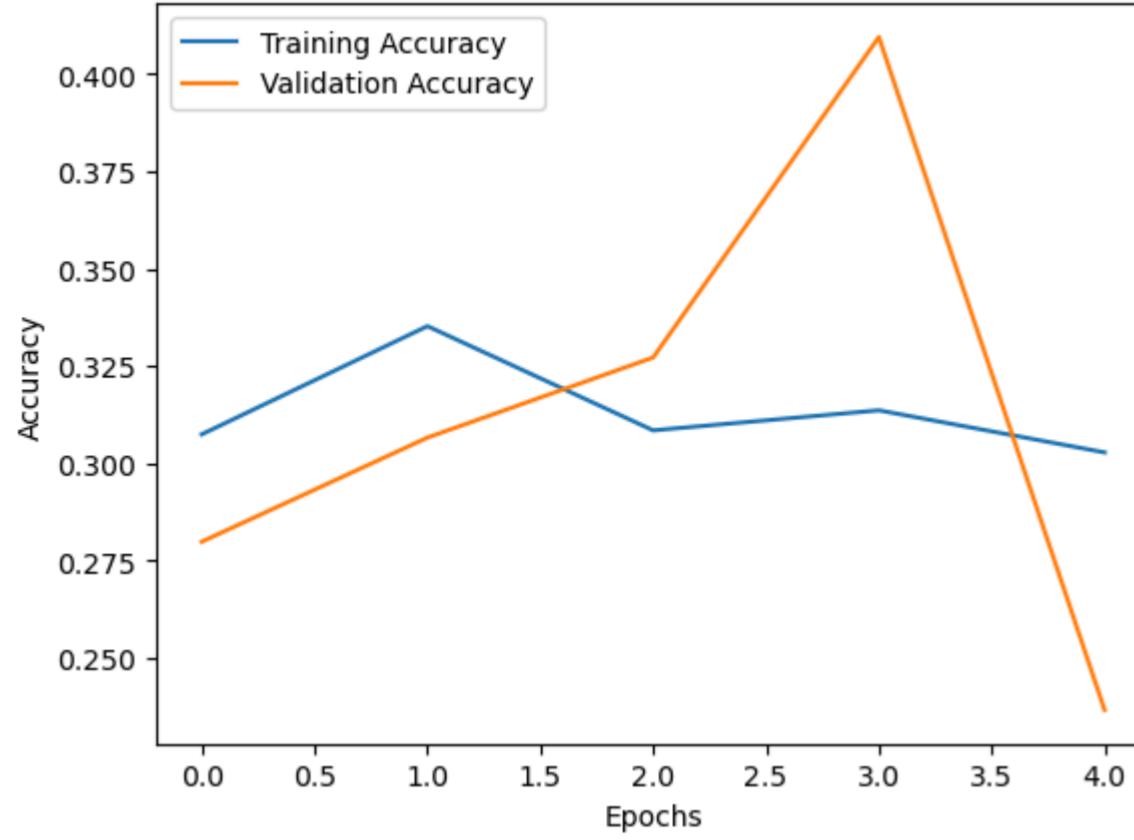
# Train the model
history = model.fit(X_train, y_train, epochs=5, batch_size=4, validation_split=0.2)

Epoch 1/5
486/486 [=====] - 5s 7ms/step - loss: -21.1100 - accuracy: 0.2765 - val_loss: -70.3031 - val_accuracy: 0.3086
Epoch 2/5
486/486 [=====] - 3s 7ms/step - loss: -115.8501 - accuracy: 0.2827 - val_loss: -226.5943 - val_accuracy: 0.2263
Epoch 3/5
486/486 [=====] - 3s 7ms/step - loss: -307.2403 - accuracy: 0.2570 - val_loss: -487.6276 - val_accuracy: 0.2716
Epoch 4/5
486/486 [=====] - 3s 6ms/step - loss: -613.3693 - accuracy: 0.2703 - val_loss: -875.3435 - val_accuracy: 0.2490
Epoch 5/5
486/486 [=====] - 3s 7ms/step - loss: -966.9393 - accuracy: 0.2585 - val_loss: -1339.9371 - val_accuracy: 0.2366
```

```
In [13]: import matplotlib.pyplot as plt

# Plotting akurasi
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Plotting kerugian
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [14]: # Evaluate the model on the testing set
predict_x = model.predict(X_test)
y_pred = np.argmax(predict_x, axis=1)

dl_acc = accuracy_score(y_pred, y_test)
dl_prec = precision_score(y_test, y_pred, average='weighted')
dl_rec = recall_score(y_test, y_pred, average='weighted')
dl_f1 = f1_score(y_test, y_pred, average='weighted')

51/51 [=====] - 0s 2ms/step
```

```
In [15]: storeResults('DNN', dl_acc, dl_prec, dl_rec, dl_f1)
```

```
In [14]: lime_explainer = lime_tabular.LimeTabularExplainer(training_data=np.array(X_train), feature_names=X_train.columns,
                                                       class_names=['0', '1', '2', '3', '4'], mode='classification')
explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=model.predict, top_labels=6, num_features=19)

157/157 [=====] - 0s 2ms/step
```

```
In [15]: explanation.show_in_notebook()
```

ML

```
In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state = 42)
```

AdaBoost

```
In [17]: from sklearn.ensemble import AdaBoostClassifier
ada = AdaBoostClassifier()

ada.fit(X_train, y_train)

y_pred = ada.predict(X_test)

ada_acc = accuracy_score(y_pred, y_test)
ada_prec = precision_score(y_pred, y_test, average='weighted')
ada_rec = recall_score(y_pred, y_test, average='weighted')
ada_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [18]: storeResults('AdaBoost', ada_acc, ada_prec, ada_rec, ada_f1)
```

LIME

```
In [18]: lime_explainer = lime_tabular.LimeTabularExplainer(training_data=np.array(X_train), feature_names=X_train.columns,
    class_names=['0', '1', '2', '3', '4'], mode='classification')
explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=ada.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

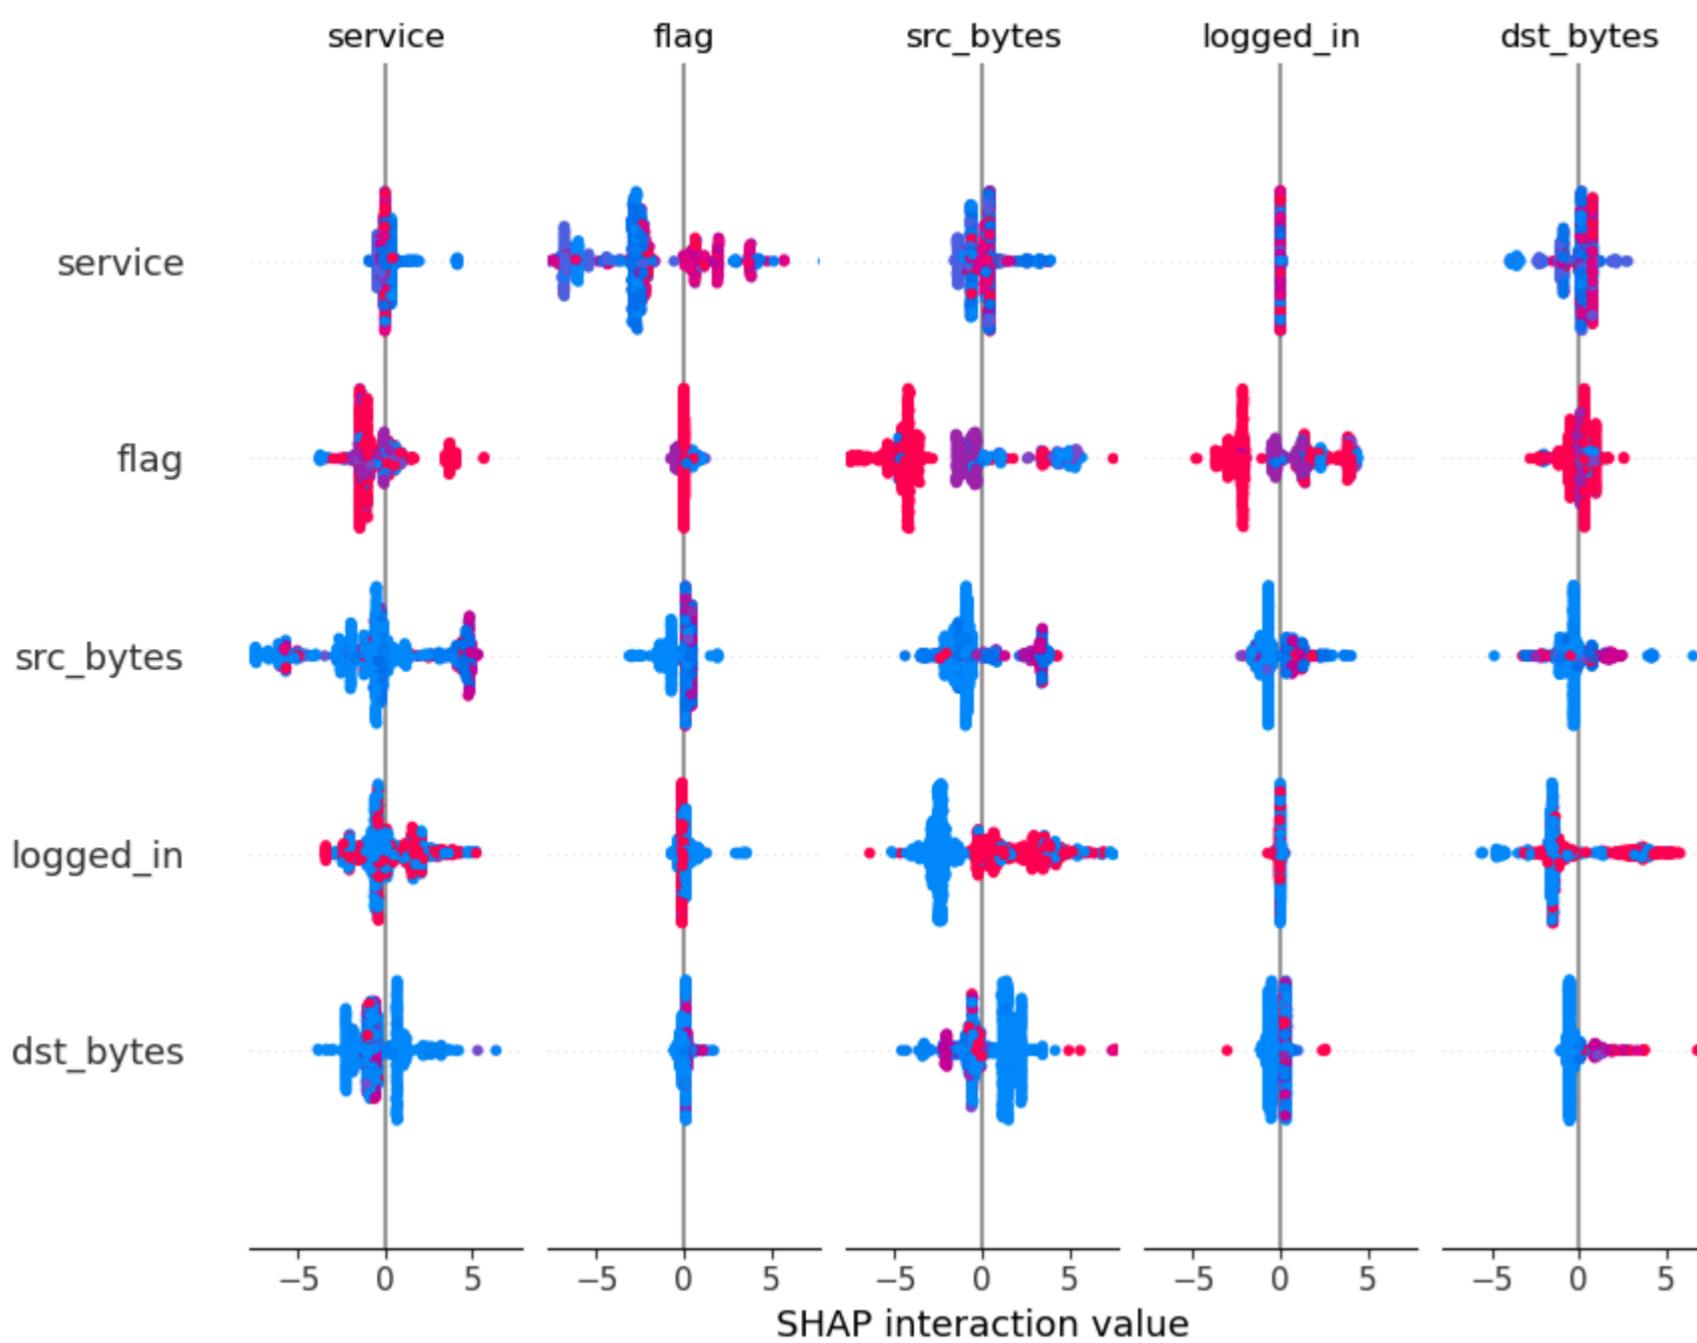
LightGBM

```
In [21]: storeResults('lightGBM',lgb.acc,lgb.prec,lgb.rec,lgb.f1)
```

LIME

```
In [21]: lime_explainer = lime_tabular.LimeTabularExplainer(training_data=np.array(X_train), feature_names=X_train.columns,
    class_names=['0', '1', '2', '3', '4'], mode='classification')
explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=clf.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

```
In [22]: explainer = shap.Explainer(clf)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```



MLP

```
In [23]: from sklearn.neural_network import MLPClassifier
mlp = MLPClassifier(random_state=1, max_iter=300)

mlp.fit(X_train, y_train)

y_pred = mlp.predict(X_test)

mlp_acc = accuracy_score(y_pred, y_test)
mlp_prec = precision_score(y_pred, y_test, average='weighted')
mlp_rec = recall_score(y_pred, y_test, average='weighted')
mlp_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [23]: storeResults('MLP', mlp_acc, mlp_prec, mlp_rec, mlp_f1)
```

```
In [24]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=mlp.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

KNN

```
In [25]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

knn_acc = accuracy_score(y_pred, y_test)
knn_prec = precision_score(y_pred, y_test, average='weighted')
knn_rec = recall_score(y_pred, y_test, average='weighted')
knn_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [25]: storeResults('KNN', knn_acc, knn_prec, knn_rec, knn_f1)
```

```
In [26]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=knn.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

Random Forest

```
In [27]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(criterion='entropy', max_features='log2', max_depth=20, n_estimators=600, min_samples_leaf=2)
rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)

rf_acc = accuracy_score(y_pred, y_test)
rf_prec = precision_score(y_pred, y_test, average='weighted')
rf_rec = recall_score(y_pred, y_test, average='weighted')
rf_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [27]: storeResults('RandomForest', rf_acc, rf_prec, rf_rec, rf_f1)
```

```
In [28]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=rf.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

SVM

```
In [29]: from sklearn import svm
svc = svm.SVC(decision_function_shape='ovo', probability=True)
svc.fit(X_train, y_train)

y_pred = svc.predict(X_test)

svc_acc = accuracy_score(y_pred, y_test)
svc_prec = precision_score(y_pred, y_test, average='weighted')
svc_rec = recall_score(y_pred, y_test, average='weighted')
svc_f1 = f1_score(y_pred, y_test, average='weighted')

In [30]: storeResults('SVM', svc_acc, svc_prec, svc_rec, svc_f1)

In [31]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=svc.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

Extension

```
In [32]: from sklearn.ensemble import VotingClassifier, BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

brf = BaggingClassifier(RandomForestClassifier())

bdt = AdaBoostClassifier(
    DecisionTreeClassifier(max_depth=1), algorithm="SAMME", n_estimators=200
)

model = VotingClassifier(estimators=[('BoostDT', bdt), ('BagRF', brf)], voting='soft')

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

ext_acc = accuracy_score(y_pred, y_test)
ext_prec = precision_score(y_pred, y_test, average='weighted')
ext_rec = recall_score(y_pred, y_test, average='weighted')
ext_f1 = f1_score(y_pred, y_test, average='weighted')

In [31]: storeResults('Extension', ext_acc, ext_prec, ext_rec, ext_f1)

In [33]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=model.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

Comparison

```
In [32]: #creating dataframe
result = pd.DataFrame({ 'ML Model' : ML_Model,
                        'Accuracy' : accuracy,
                        'Precision': precision,
                        'Recall' : recall,
                        'F1_score' : f1score
                      })

In [33]: result
```

	ML Model	Accuracy	Precision	Recall	F1_score
0	DNN	0.253	0.064	0.253	0.102
1	AdaBoost	0.699	0.830	0.699	0.733
2	LightGBM	0.984	0.985	0.984	0.984
3	MLP	0.847	0.837	0.847	0.840
4	KNN	0.935	0.938	0.935	0.936
5	RandomForest	0.977	0.978	0.977	0.977
6	SVM	0.257	0.984	0.257	0.387
7	Extension	0.994	0.994	0.994	0.994

Modelling

```
In [34]: import joblib
filename = 'models/model_nsl.sav'
joblib.dump(model, filename)
```

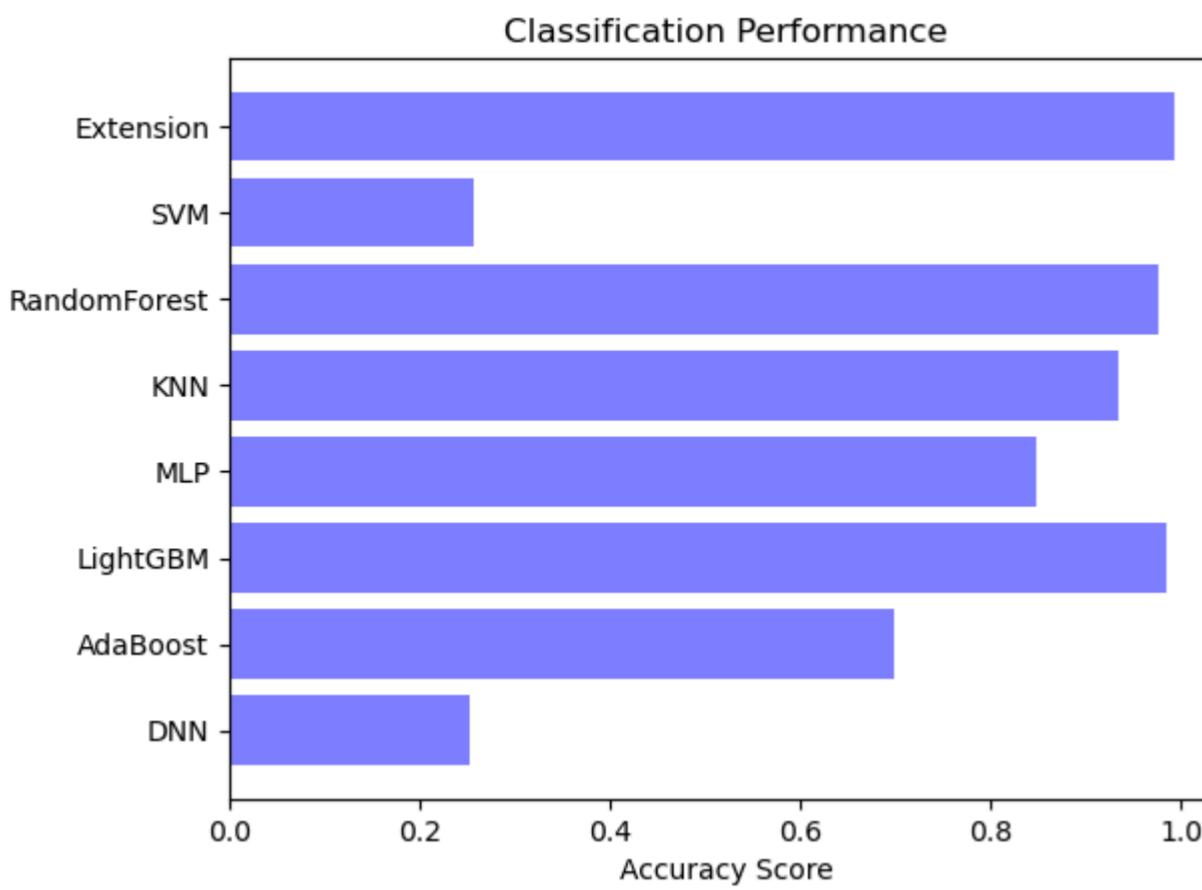
```
Out[34]: ['models/model_nsl.sav']
```

Graph

```
In [36]: classifier = ML_Model
y_pos = np.arange(len(classifier))
```

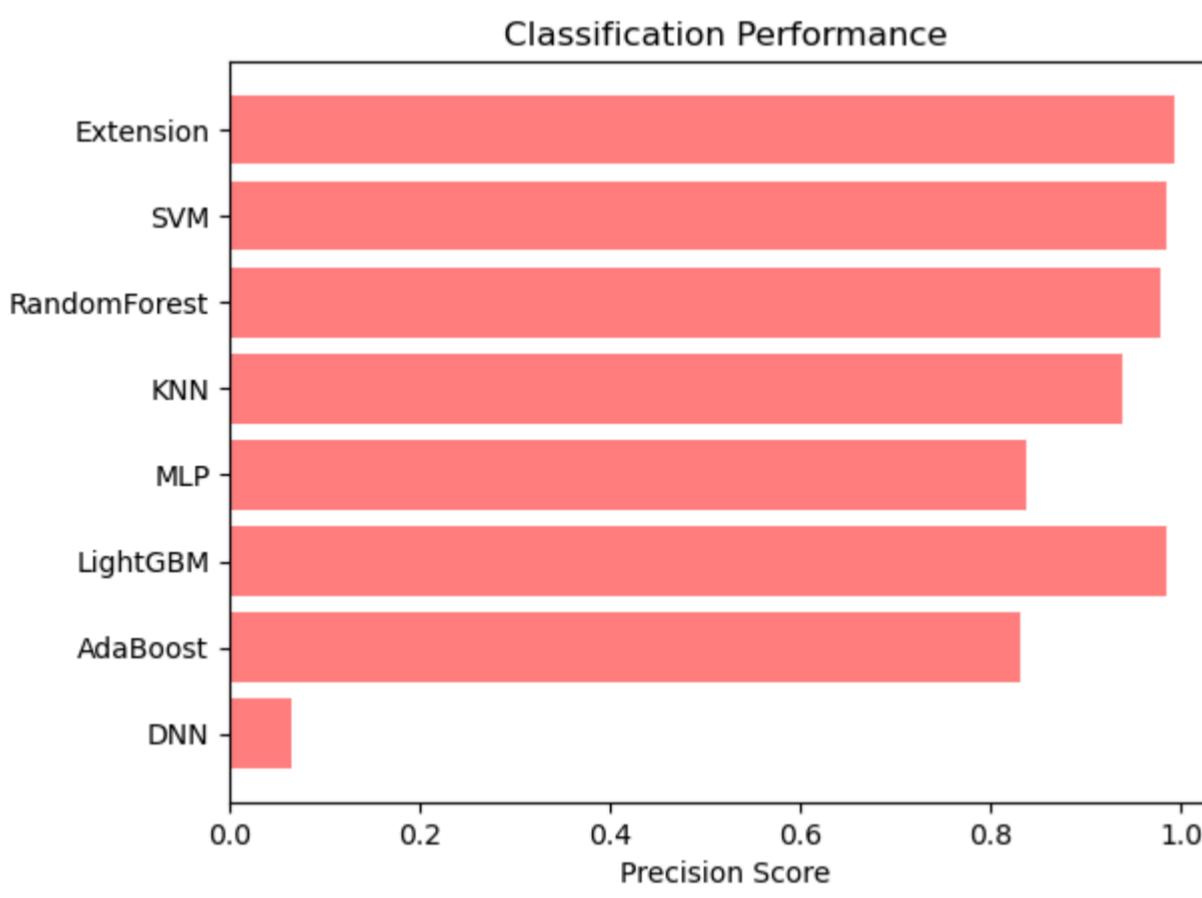
Accuracy

```
In [37]: import matplotlib.pyplot as plt2
plt2.barh(y_pos, accuracy, align='center', alpha=0.5,color='blue')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Accuracy Score')
plt2.title('Classification Performance')
plt2.show()
```



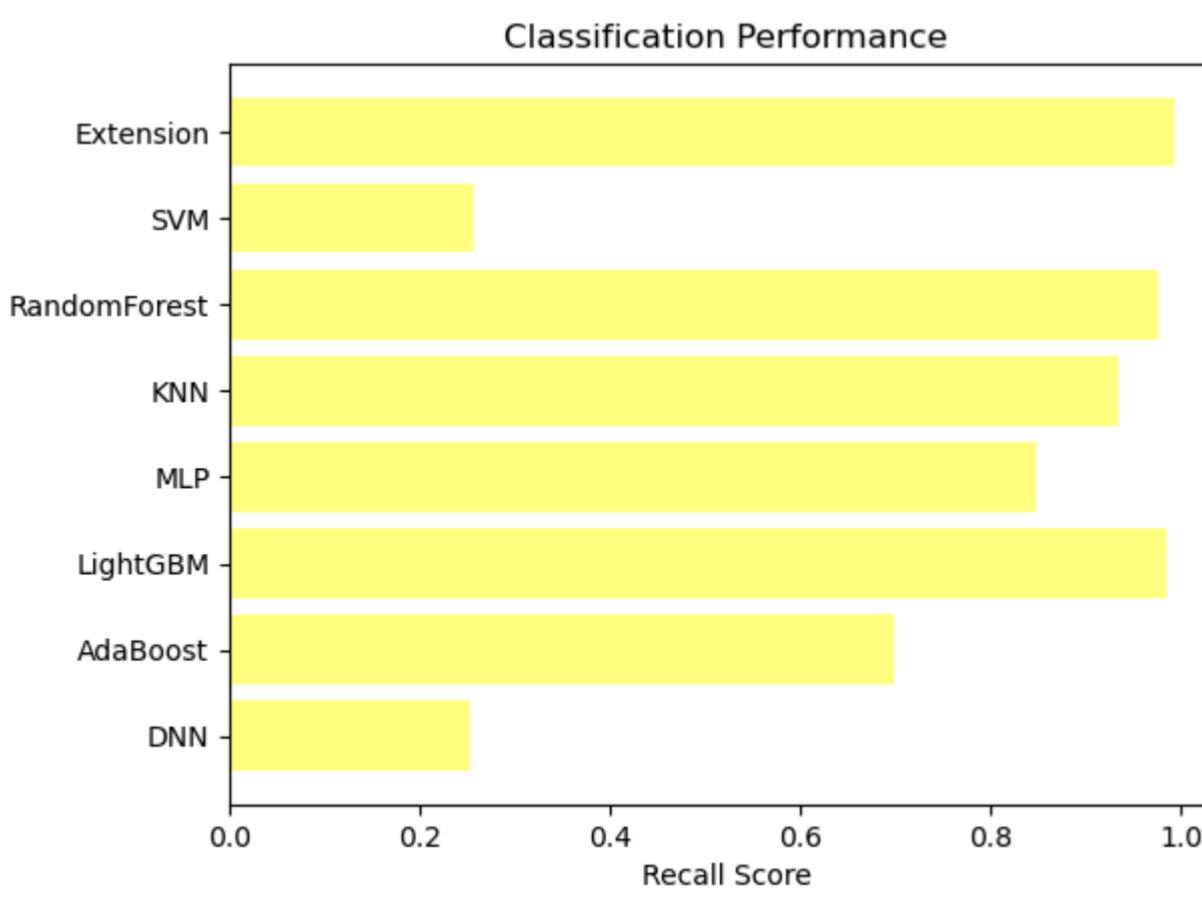
Precision

```
In [38]: plt2.barh(y_pos, precision, align='center', alpha=0.5,color='red')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Precision Score')
plt2.title('Classification Performance')
plt2.show()
```



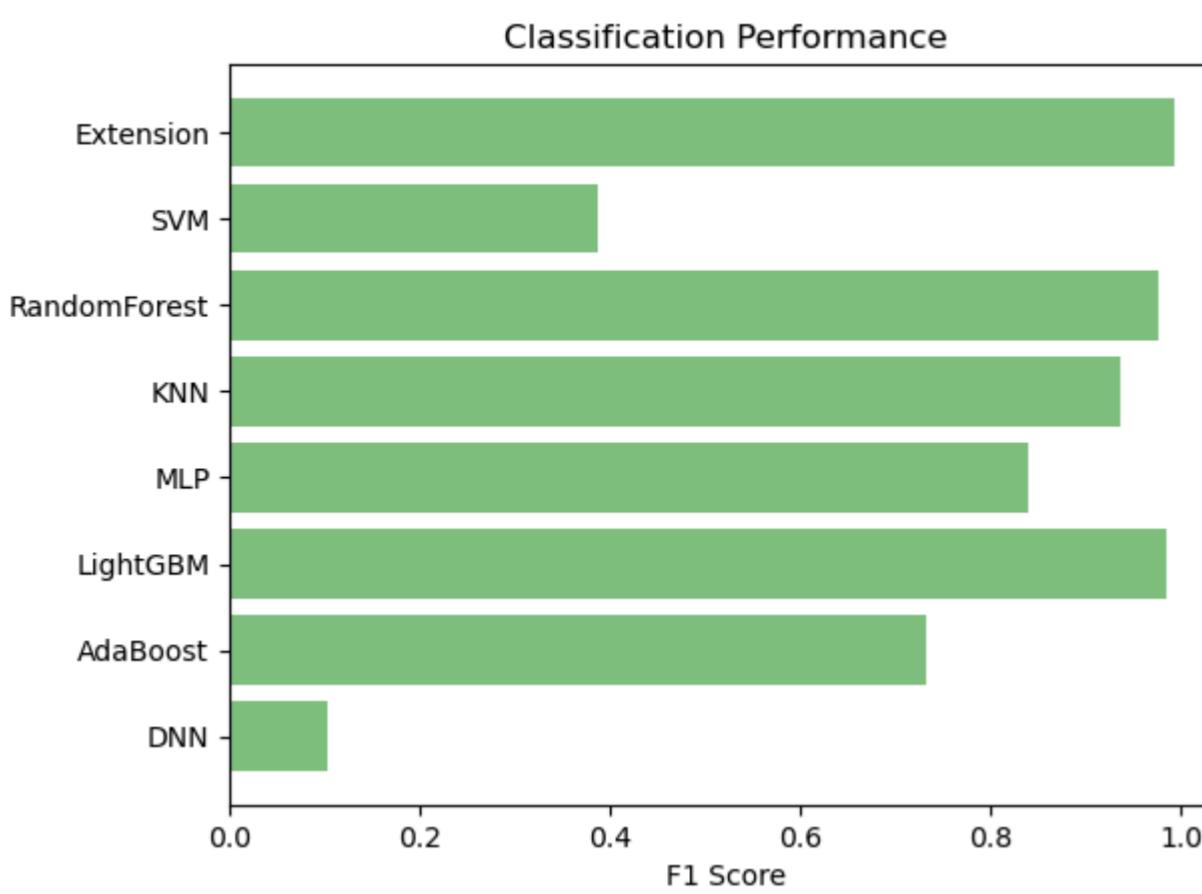
Recall

```
In [39]: plt2.barh(y_pos, recall, align='center', alpha=0.5,color='yellow')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Recall Score')
plt2.title('Classification Performance')
plt2.show()
```



F1 Score

```
In [40]: plt2.barh(y_pos, f1score, align='center', alpha=0.5,color='green')
plt2.yticks(y_pos, classifier)
plt2.xlabel('F1 Score')
plt2.title('Classification Performance')
plt2.show()
```



```
In [ ]:
```