

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [4]: data = pd.read_csv('SIMARGL2021/dataset-part1.csv')
```

```
In [5]: # Load data
data.columns
```

```
Out[5]: Index(['BIFLOW_DIRECTION', 'DIRECTION', 'DST_TO_SRC_SECOND_BYTES',
   'FIREWALL_EVENT', 'FIRST_SWITCHED', 'FLOW_ACTIVE_TIMEOUT',
   'FLOW_DURATION_MICROSECONDS', 'FLOW_DURATION_MILLISECONDS',
   'FLOW_END_MILLISECONDS', 'FLOW_END_SEC', 'FLOW_ID',
   'FLOW_INACTIVE_TIMEOUT', 'FLOW_START_MILLISECONDS', 'FLOW_START_SEC',
   'FRAME_LENGTH', 'IN_BYTES', 'IN_PKTS', 'IPV4_DST_ADDR', 'IPV4_SRC_ADDR',
   'L4_DST_PORT', 'L4_SRC_PORT', 'LAST_SWITCHED', 'MAX_IP_PKT_LEN',
   'MIN_IP_PKT_LEN', 'OOORDER_IN_PKTS', 'OOORDER_OUT_PKTS', 'OUT_BYTES',
   'OUT_PKTS', 'PROTOCOL', 'PROTOCOL_MAP', 'RETRANSMITTED_IN_BYTES',
   'RETRANSMITTED_IN_PKTS', 'RETRANSMITTED_OUT_BYTES',
   'RETRANSMITTED_OUT_PKTS', 'SRC_TO_DST_SECOND_BYTES', 'TCP_FLAGS',
   'TCP_WIN_MAX_IN', 'TCP_WIN_MAX_OUT', 'TCP_WIN_MIN_IN',
   'TCP_WIN_MIN_OUT', 'TCP_WIN_MSS_IN', 'TCP_WIN_MSS_OUT',
   'TCP_WIN_SCALE_IN', 'TCP_WIN_SCALE_OUT', 'SRC_TOS', 'DST_TOS',
   'L7_PROTO_NAME', 'SAMPLING_INTERVAL', 'TOTAL_FLOWS_EXP', 'LABEL'],
  dtype='object')
```

```
In [6]: null_counts = data.isnull().sum()
```

```
# Print the number of null values
```

```
print(f'{null_counts.sum()} null entries have been found in the dataset\n')
```

```
# Drop null values
```

```
data.dropna(inplace=True)      # or df_data = df_data.dropna()
```

```
# Find and handle duplicates
```

```
duplicate_count = data.duplicated().sum()
```

```
# Print the number of duplicate entries
```

```
print(f'{duplicate_count} duplicate entries have been found in the dataset\n')
```

```
# Remove duplicates
```

```
data.drop_duplicates(inplace=True) # or df_data = df_data.drop_duplicates()
```

```
# Display relative message
```

```
print(f'All duplicates have been removed\n')
```

```
# Reset the indexes
```

```
data.reset_index(drop=True, inplace=True)
```

```
# Inspect the dataset for categorical columns
```

```
print("Categorical columns:", data.select_dtypes(include=['object']).columns.tolist(), '\n')
```

```
# Print the first 5 lines
```

```
data.head()
```

0 null entries have been found in the dataset

0 duplicate entries have been found in the dataset

All duplicates have been removed

Categorical columns: ['DST\_TO\_SRC\_SECOND\_BYTES', 'IPV4\_DST\_ADDR', 'IPV4\_SRC\_ADDR', 'PROTOCOL\_MAP', 'SRC\_TO\_DST\_SECOND\_BYTES', 'L7\_PROTO\_NAME', 'LABEL']

	BIFLOW_DIRECTION	DIRECTION	DST_TO_SRC_SECOND_BYTES	FIREWALL_EVENT	FIRST_SWITCHED	FLOW_ACTIVE_TIMEOUT	FLOW_DURATION_MICROSECONDS	FLOW_DURATION_MILLISECONDS	FLOW_END_MILLIS
0	1	0		40	0	1616660040	120	339	0
1	1	0		,	0	1616660040	120	0	0
2	1	0		104	0	1616660040	120	44725	44
3	1	0		,	0	1616660040	120	0	0
4	1	0		40	0	1616660040	120	1114	1

5 rows × 50 columns

```
In [7]: drop_columns = [ # this list includes all spellings across CIC NIDS datasets
    'DST_TO_SRC_SECOND_BYTES', 'IPV4_DST_ADDR', 'IPV4_SRC_ADDR', 'PROTOCOL_MAP', 'SRC_TO_DST_SECOND_BYTES', 'L7_PROTO_NAME'
]
data.drop(columns=drop_columns, inplace=True, errors='ignore')
```

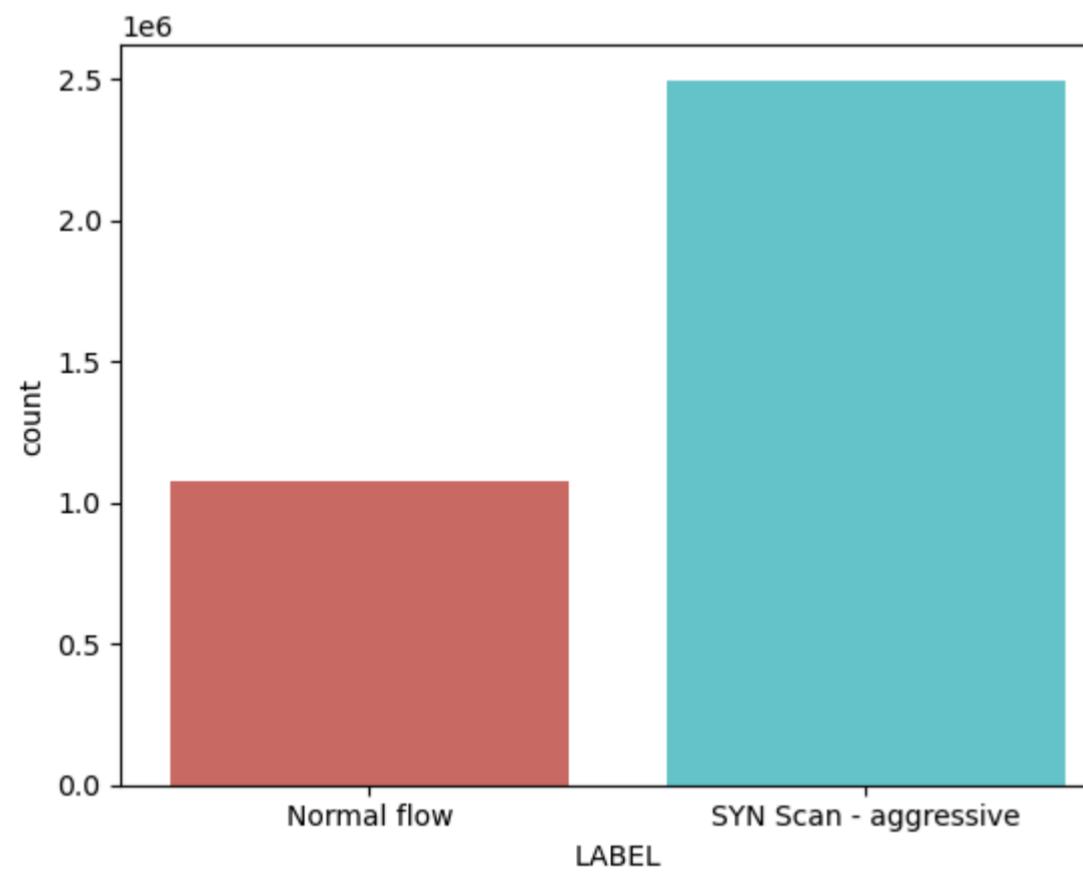
In [8]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3570666 entries, 0 to 3570665
Data columns (total 44 columns):
 #   Column           Dtype  
 --- 
 0   BIFLOW_DIRECTION    int64  
 1   DIRECTION          int64  
 2   FIREWALL_EVENT     int64  
 3   FIRST_SWITCHED     int64  
 4   FLOW_ACTIVE_TIMEOUT int64  
 5   FLOW_DURATION_MICROSECONDS int64  
 6   FLOW_DURATION_MILLISECONDS int64  
 7   FLOW_END_MILLISECONDS int64  
 8   FLOW_END_SEC        int64  
 9   FLOW_ID             int64  
 10  FLOW_INACTIVE_TIMEOUT int64  
 11  FLOW_START_MILLISECONDS int64  
 12  FLOW_START_SEC      int64  
 13  FRAME_LENGTH        int64  
 14  IN_BYTES            int64  
 15  IN_PKTS             int64  
 16  L4_DST_PORT         int64  
 17  L4_SRC_PORT          int64  
 18  LAST_SWITCHED       int64  
 19  MAX_IP_PKT_LEN      int64  
 20  MIN_IP_PKT_LEN      int64  
 21  OOORDER_IN_PKTS     int64  
 22  OOORDER_OUT_PKTS    int64  
 23  OUT_BYTES            int64  
 24  OUT_PKTS             int64  
 25  PROTOCOL             int64  
 26  RETRANSMITTED_IN_BYTES int64  
 27  RETRANSMITTED_IN_PKTS int64  
 28  RETRANSMITTED_OUT_BYTES int64  
 29  RETRANSMITTED_OUT_PKTS int64  
 30  TCP_FLAGS             int64  
 31  TCP_WIN_MAX_IN       int64  
 32  TCP_WIN_MAX_OUT      int64  
 33  TCP_WIN_MIN_IN       int64  
 34  TCP_WIN_MIN_OUT      int64  
 35  TCP_WIN_MSS_IN       int64  
 36  TCP_WIN_MSS_OUT      int64  
 37  TCP_WIN_SCALE_IN     int64  
 38  TCP_WIN_SCALE_OUT    int64  
 39  SRC_TOS              int64  
 40  DST_TOS              int64  
 41  SAMPLING_INTERVAL    int64  
 42  TOTAL_FLOWS_EXP     int64  
 43  LABEL                object 
dtypes: int64(43), object(1)
memory usage: 1.2+ GB
```

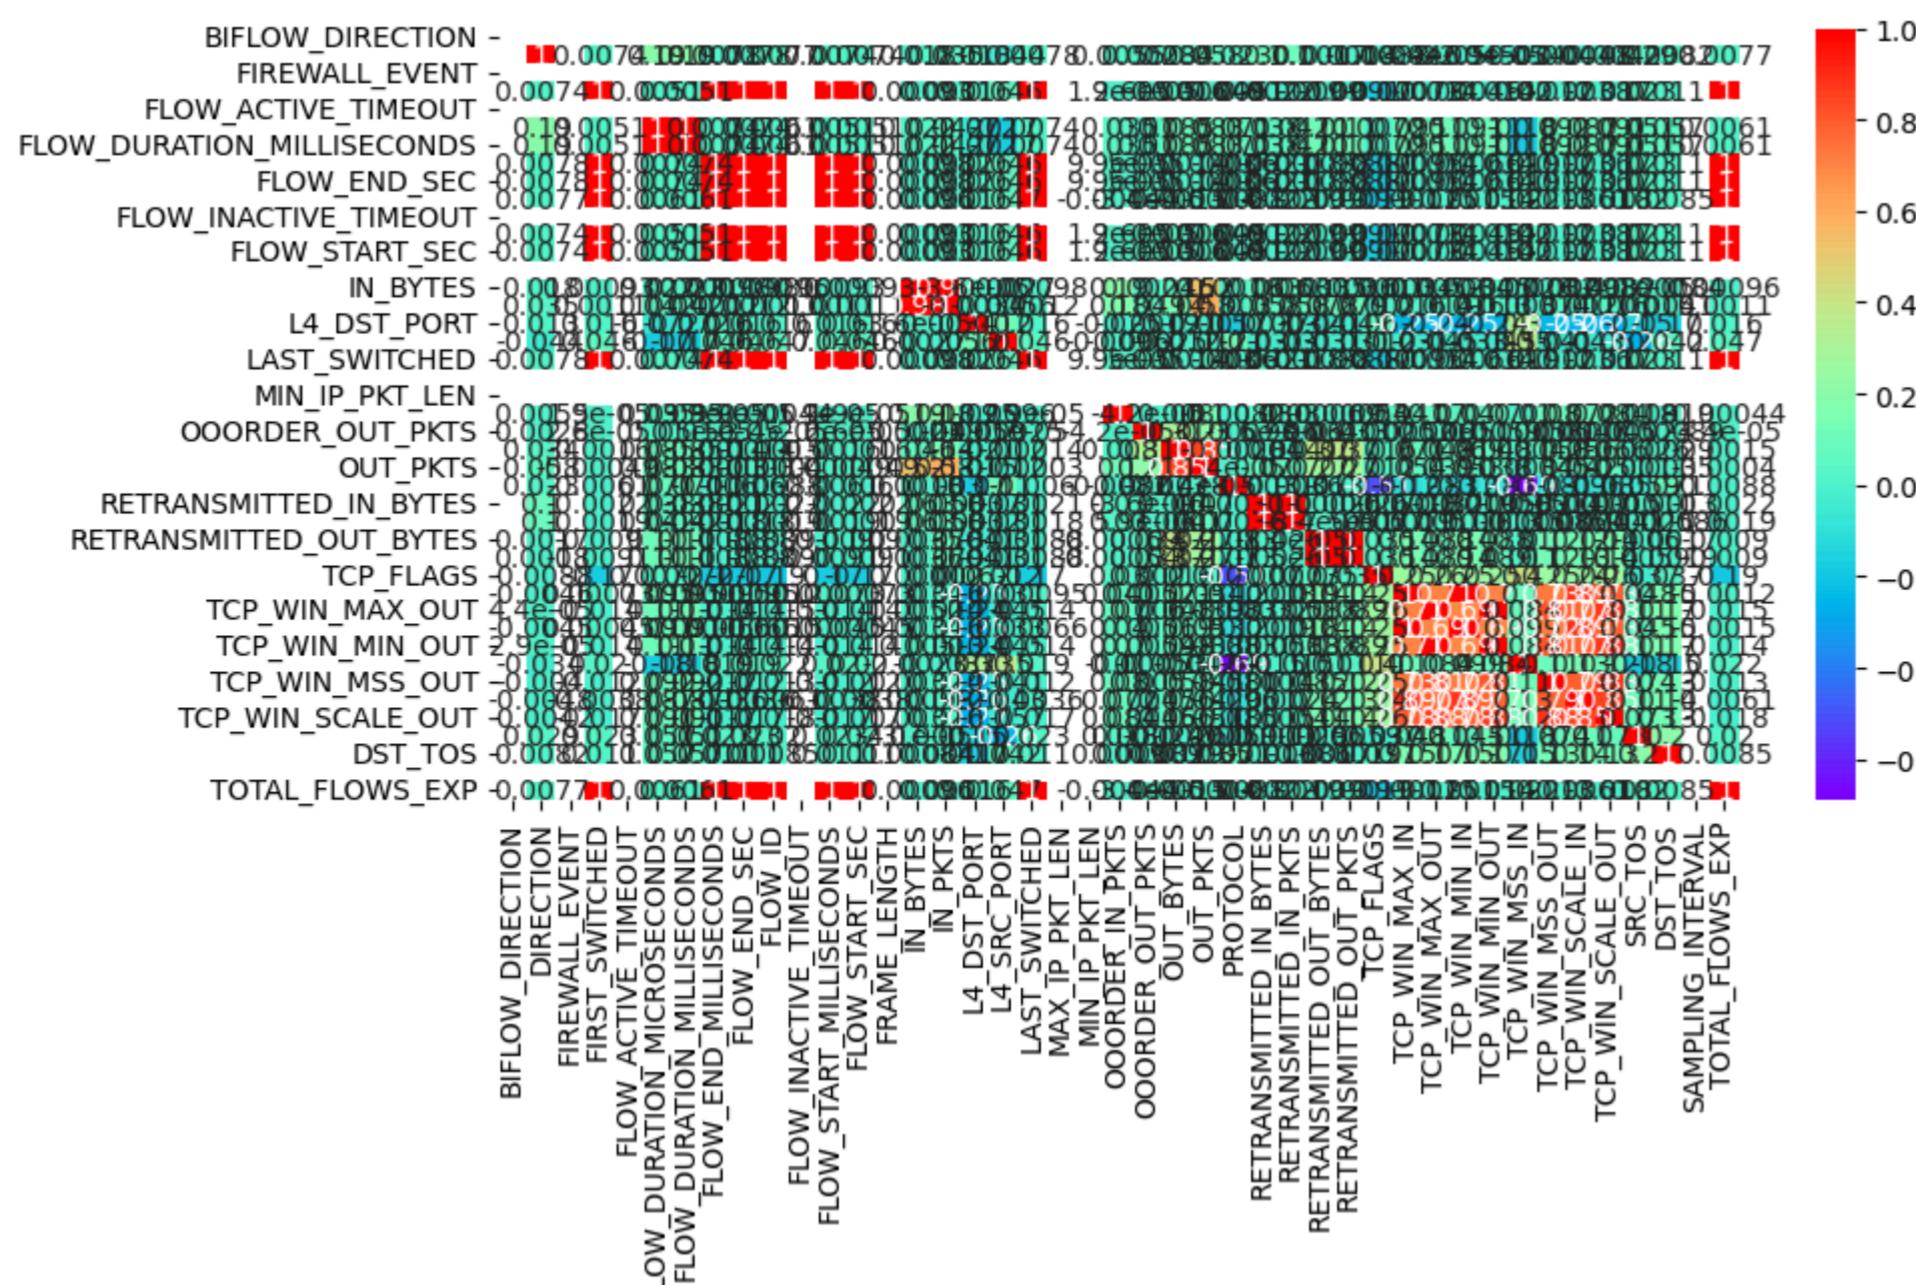
In [9]: `data['LABEL'].value_counts()`

```
Out[9]: SYN Scan - aggressive    2496814
Normal flow                  1073852
Name: LABEL, dtype: int64
```

```
In [10]: sns.countplot(x='LABEL', data=data, palette='hls')
plt.show()
#plt.savefig('count_plot') mal: the nodule malignancy, 0: benign, 1: malignant
```



```
In [11]: plt.figure(figsize = (10,5))
sns.heatmap(data.corr(), annot = True, cmap="rainbow")
plt.show()
```



```
In [14]: # Import Label encoder
from sklearn import preprocessing

# Label_encoder object knows
# how to understand word labels.
label_encoder = preprocessing.LabelEncoder()

# Encode Labels in column 'species'.
data['LABEL']= label_encoder.fit_transform(data['LABEL'])

data['LABEL'].unique()
```

Out[14]: array([1, 0])

```
In [15]: X = data.drop(["LABEL"],axis =1)
y = data["LABEL"]
```

## FS

```
In [16]: from sklearn.feature_selection import SelectKBest, SelectPercentile, mutual_info_classif
```

```
In [17]: selector = SelectPercentile(mutual_info_classif, percentile=25)
X_reduced = selector.fit_transform(X, y)
#X_reduced.shape
```

```
In [19]: cols = selector.get_support(indices=True)
selected_columns = X.iloc[:,cols].columns.tolist()
selected_columns
```

```
Out[19]: ['FLOW_DURATION_MICROSECONDS',
'FLOW_DURATION_MILLISECONDS',
'IN_BYTES',
'L4_DST_PORT',
'L4_SRC_PORT',
'OUT_BYTES',
'PROTOCOL',
'TCP_FLAGS',
'TCP_WIN_MAX_IN',
'TCP_WIN_MIN_IN',
'TCP_WIN_MSS_IN']
```

```
In [20]: len(selected_columns)
```

Out[20]: 11

```
In [21]: df = data[['FLOW_DURATION_MICROSECONDS',
'FLOW_DURATION_MILLISECONDS',
'IN_BYTES',
'L4_DST_PORT',
'L4_SRC_PORT',
'OUT_BYTES',
'PROTOCOL',
'TCP_FLAGS',
'TCP_WIN_MAX_IN',
'TCP_WIN_MIN_IN',
'TCP_WIN_MSS_IN',
'LABEL']]
```

```
In [22]: df.to_csv('simarg_processed.csv')
```

```
In [3]: df = pd.read_csv('simarg_processed.csv')
```

```
In [4]: del df['Unnamed: 0']
```

```
In [5]: df.columns
```

```
Out[5]: Index(['FLOW_DURATION_MICROSECONDS', 'FLOW_DURATION_MILLISECONDS', 'IN_BYTES',
'L4_DST_PORT', 'L4_SRC_PORT', 'OUT_BYTES', 'PROTOCOL', 'TCP_FLAGS',
'TCP_WIN_MAX_IN', 'TCP_WIN_MIN_IN', 'TCP_WIN_MSS_IN', 'LABEL'],
dtype='object')
```

```
In [6]: X = df[['FLOW_DURATION_MICROSECONDS', 'FLOW_DURATION_MILLISECONDS', 'IN_BYTES',
   'L4_DST_PORT', 'L4_SRC_PORT', 'OUT_BYTES', 'PROTOCOL', 'TCP_FLAGS',
   'TCP_WIN_MAX_IN', 'TCP_WIN_MIN_IN', 'TCP_WIN_MSS_IN']]
y = df["LABEL"]
```

```
In [7]: from sklearn.model_selection import train_test_split
```

```
In [8]: from sklearn.metrics import accuracy_score # for calculating accuracy of model
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
In [9]: ML_Model = []
accuracy = []
precision = []
recall = []

f1score = []

#function to call for storing the results
def storeResults(model, a,b,c,d):
    ML_Model.append(model)
    accuracy.append(round(a, 3))
    precision.append(round(b, 3))
    recall.append(round(c, 3))
    f1score.append(round(d, 3))
```

```
In [15]: # importing lime
import lime
from lime import lime_tabular
import shap
```

## DNN

```
In [10]: from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)
# Scale the features using StandardScaler
#scaler = StandardScaler()
#X_train = scaler.fit_transform(X_train)
#X_test = scaler.transform(X_test)
```

```
In [13]: # Build the model architecture
model = tf.keras.Sequential([
    tf.keras.layers.Dense(1024, activation='relu', input_dim=X_train.shape[1]),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(256, activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(128, activation='tanh'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])

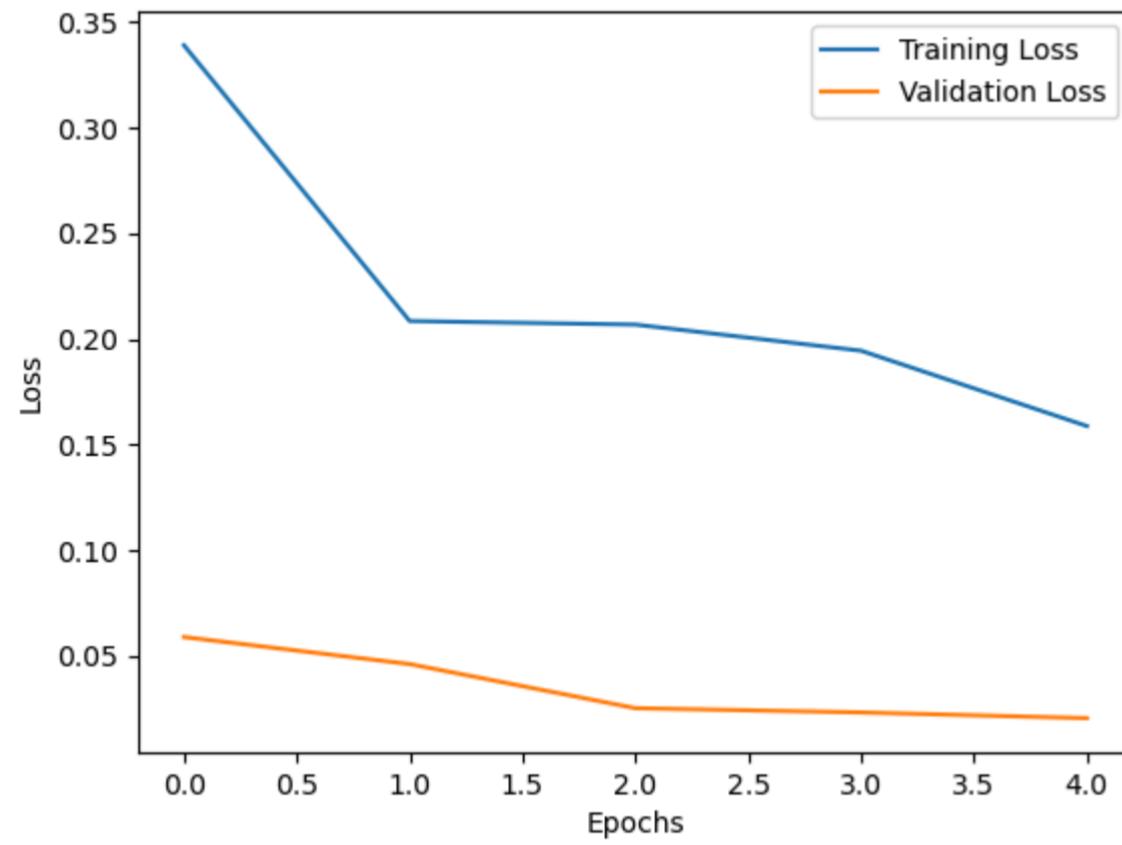
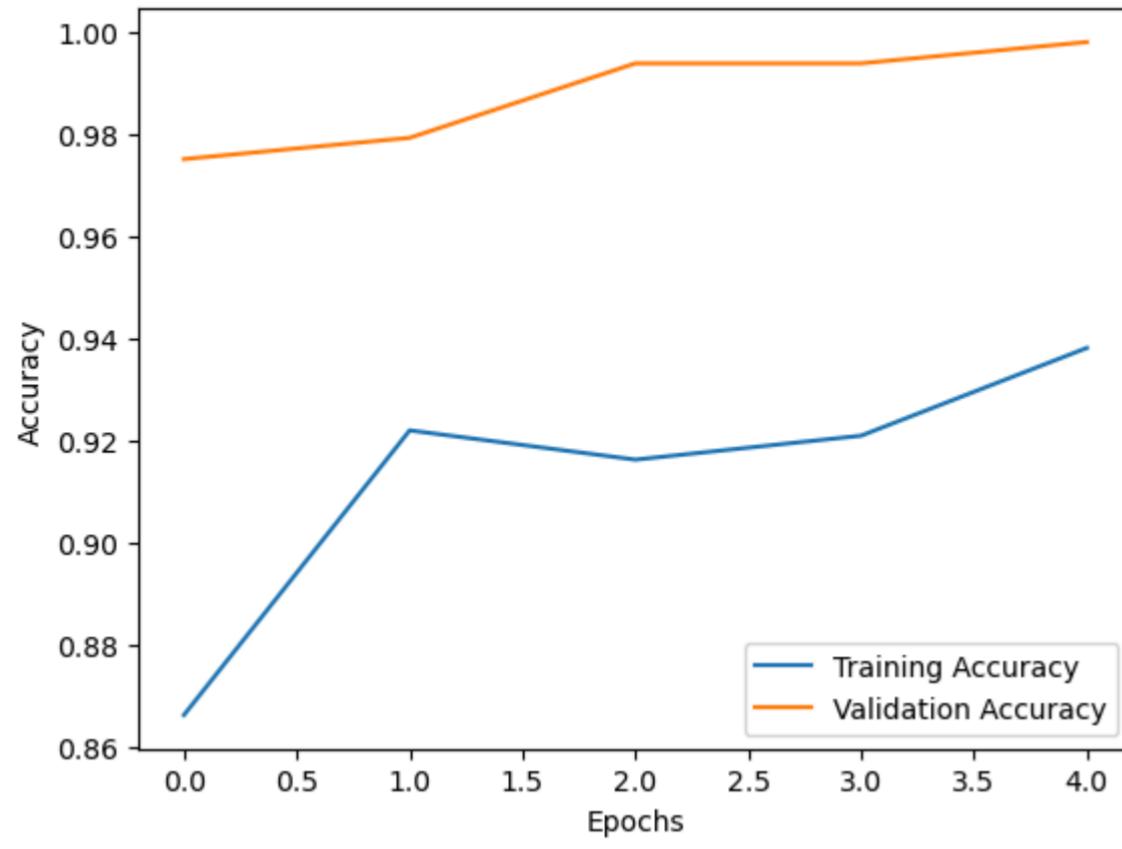
# Train the model
history = model.fit(X_train, y_train, epochs=5, batch_size=4, validation_split=0.2)

Epoch 1/5
480/480 [=====] - 5s 6ms/step - loss: 0.6289 - accuracy: 0.7318 - val_loss: 0.3588 - val_accuracy: 0.8583
Epoch 2/5
480/480 [=====] - 2s 5ms/step - loss: 0.5105 - accuracy: 0.7729 - val_loss: 0.3445 - val_accuracy: 0.8792
Epoch 3/5
480/480 [=====] - 2s 5ms/step - loss: 0.4602 - accuracy: 0.7995 - val_loss: 0.4112 - val_accuracy: 0.8104
Epoch 4/5
480/480 [=====] - 2s 5ms/step - loss: 0.4237 - accuracy: 0.8245 - val_loss: 0.3230 - val_accuracy: 0.8813
Epoch 5/5
480/480 [=====] - 2s 5ms/step - loss: 0.3814 - accuracy: 0.8448 - val_loss: 0.2907 - val_accuracy: 0.9688
```

```
In [13]: import matplotlib.pyplot as plt

# Plotting akurasi
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Plotting kerugian
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [14]: # Evaluate the model on the testing set
predict_x = model.predict(X_test)
y_pred = np.argmax(predict_x, axis=1)

dl_acc = accuracy_score(y_pred, y_test)
dl_prec = precision_score(y_test, y_pred, average='weighted')
dl_rec = recall_score(y_test, y_pred, average='weighted')
dl_f1 = f1_score(y_test, y_pred, average='weighted')

50/50 [=====] - 0s 2ms/step
```

```
In [15]: storeResults('DNN', dl_acc, dl_prec, dl_rec, dl_f1)
```

```
In [16]: lime_explainer = lime_tabular.LimeTabularExplainer(training_data=np.array(X_train), feature_names=X_train.columns,
                                                       class_names=['0', '1'], mode='classification')
explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=model.predict, top_labels=6, num_features=19)

157/157 [=====] - 0s 2ms/step
```

```
In [17]: explanation.show_in_notebook()
```

ML

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state = 42)
```

## AdaBoost

```
In [19]: from sklearn.ensemble import AdaBoostClassifier
ada = AdaBoostClassifier()

ada.fit(X_train, y_train)

y_pred = ada.predict(X_test)

ada_acc = accuracy_score(y_pred, y_test)
ada_prec = precision_score(y_pred, y_test, average='weighted')
ada_rec = recall_score(y_pred, y_test, average='weighted')
ada_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [18]: storeResults('AdaBoost', ada_acc, ada_prec, ada_rec, ada_f1)
```

```
In [20]: lime_explainer = lime_tabular.LimeTabularExplainer(training_data=np.array(X_train), feature_names=X_train.columns,
                                                       class_names=['0','1'], mode='classification')
explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=ada.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

## LightGBM

```
In [22]: # build the lightgbm model
import lightgbm as lgb
clf = lgb.LGBMClassifier()
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

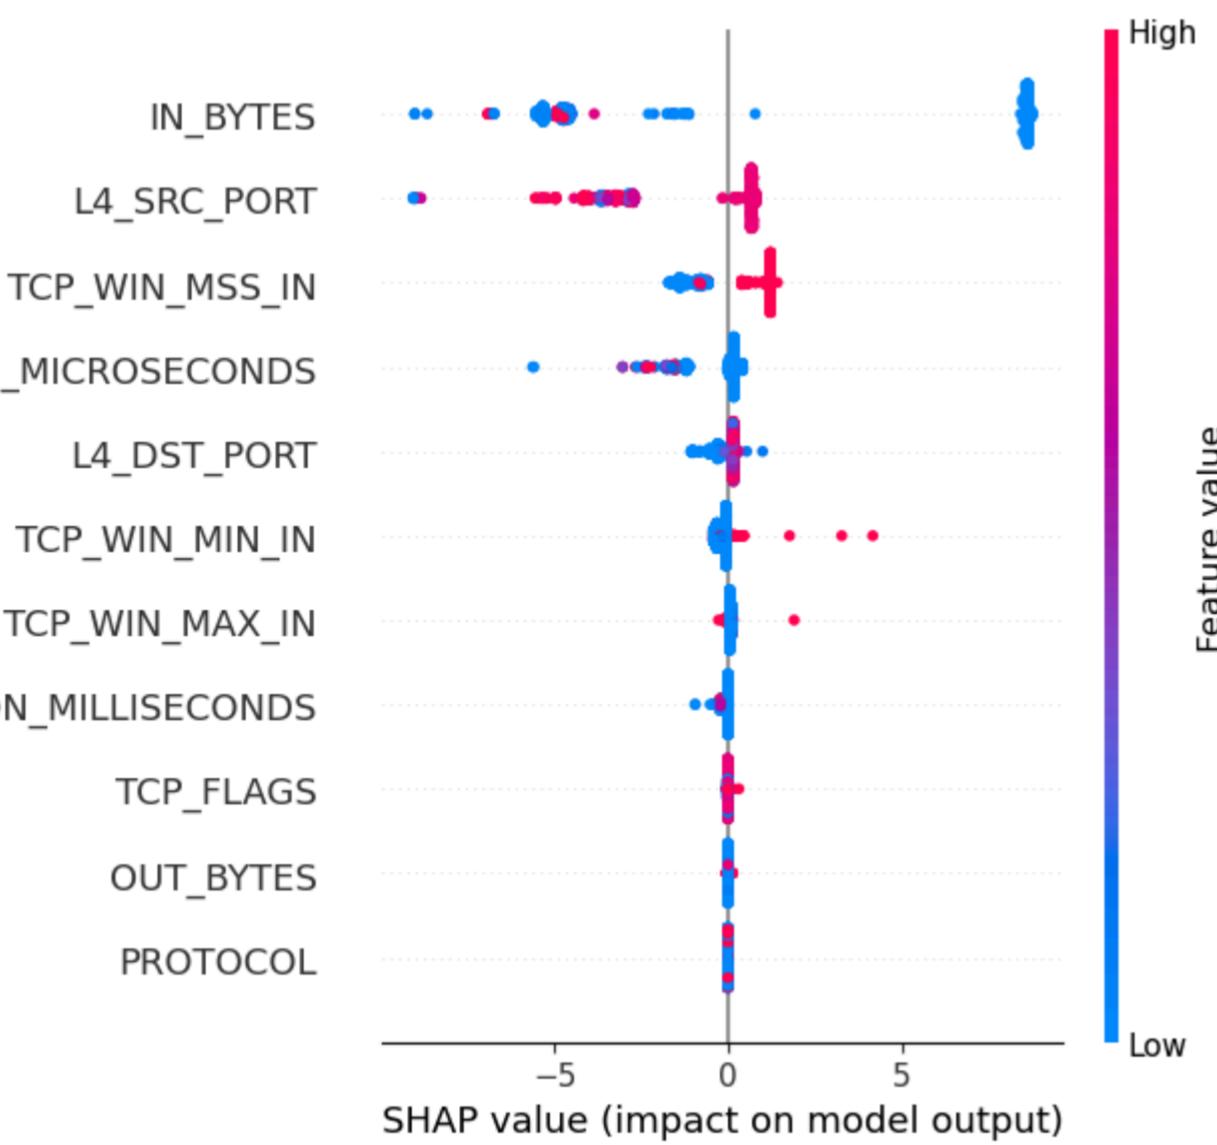
lgb_acc = accuracy_score(y_pred, y_test)
lgb_prec = precision_score(y_pred, y_test, average='weighted')
lgb_rec = recall_score(y_pred, y_test, average='weighted')
lgb_f1 = f1_score(y_pred, y_test, average='weighted')

[LightGBM] [Info] Number of positive: 1578, number of negative: 1622
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000625 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 1698
[LightGBM] [Info] Number of data points in the train set: 3200, number of used features: 11
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.493125 -> initscore=-0.027502
[LightGBM] [Info] Start training from score -0.027502
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

```
In [21]: storeResults('LightGBM',lgb_acc,lgb_prec,lgb_rec,lgb_f1)
```

```
In [23]: lime_explainer = lime_tabular.LimeTabularExplainer(training_data=np.array(X_train), feature_names=X_train.columns,
                                                       class_names=['0','1'], mode='classification')
explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=clf.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

```
In [24]: explainer = shap.Explainer(clf)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```



## MLP

```
In [25]: from sklearn.neural_network import MLPClassifier
mlp = MLPClassifier(random_state=1, max_iter=300)

mlp.fit(X_train, y_train)

y_pred = mlp.predict(X_test)

mlp_acc = accuracy_score(y_pred, y_test)
mlp_prec = precision_score(y_pred, y_test, average='weighted')
mlp_rec = recall_score(y_pred, y_test, average='weighted')
mlp_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [23]: storeResults('MLP',mlp_acc,mlp_prec,mlp_rec,mlp_f1)
```

```
In [26]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=mlp.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

## KNN

```
In [27]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

knn_acc = accuracy_score(y_pred, y_test)
knn_prec = precision_score(y_pred, y_test, average='weighted')
knn_rec = recall_score(y_pred, y_test, average='weighted')
knn_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [25]: storeResults('KNN', knn_acc, knn_prec, knn_rec, knn_f1)
```

```
In [28]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=knn.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

## Random Forest

```
In [29]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(criterion='entropy', max_features='log2', max_depth=20, n_estimators=600, min_samples_leaf=2)
rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)

rf_acc = accuracy_score(y_pred, y_test)
rf_prec = precision_score(y_pred, y_test, average='weighted')
rf_rec = recall_score(y_pred, y_test, average='weighted')
rf_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [27]: storeResults('RandomForest', rf_acc, rf_prec, rf_rec, rf_f1)
```

```
In [30]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=rf.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

## SVM

```
In [31]: from sklearn import svm
svc = svm.SVC(decision_function_shape='ovo', probability=True)
svc.fit(X_train, y_train)

y_pred = svc.predict(X_test)

svc_acc = accuracy_score(y_pred, y_test)
svc_prec = precision_score(y_pred, y_test, average='weighted')
svc_rec = recall_score(y_pred, y_test, average='weighted')
svc_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [29]: storeResults('SVM', svc_acc, svc_prec, svc_rec, svc_f1)
```

```
In [32]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=svc.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

## Extension

```
In [33]: from sklearn.ensemble import VotingClassifier, BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

brf = BaggingClassifier(RandomForestClassifier())

bdt = AdaBoostClassifier(
    DecisionTreeClassifier(max_depth=1), algorithm="SAMME", n_estimators=200
)

model = VotingClassifier(estimators=[('BoostDT', bdt), ('BagRF', brf)], voting='soft')

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

ext_acc = accuracy_score(y_pred, y_test)
ext_prec = precision_score(y_pred, y_test, average='weighted')
ext_rec = recall_score(y_pred, y_test, average='weighted')
ext_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [31]: storeResults('Extension', ext_acc, ext_prec, ext_rec, ext_f1)
```

```
In [34]: explanation = lime_explainer.explain_instance(data_row=X_test.iloc[1], predict_fn=model.predict_proba, top_labels=6, num_features=19)
explanation.show_in_notebook()
```

## Comparison

```
In [32]: #creating dataframe
result = pd.DataFrame({
    'ML Model': ML_Model,
    'Accuracy': accuracy,
    'Precision': precision,
    'Recall': recall,
    'F1 score': f1score
})
```

In [33]: result

	ML Model	Accuracy	Precision	Recall	F1_score
0	DNN	0.487	0.237	0.487	0.319
1	AdaBoost	0.999	0.999	0.999	0.999
2	LightGBM	0.999	0.999	0.999	0.999
3	MLP	1.000	1.000	1.000	1.000
4	KNN	0.998	0.998	0.998	0.998
5	RandomForest	1.000	1.000	1.000	1.000
6	SVM	0.628	0.921	0.628	0.700
7	Extension	1.000	1.000	1.000	1.000

## Modelling

```
In [34]: import joblib
filename = 'models/model_simarg.sav'
joblib.dump(model, filename)
```

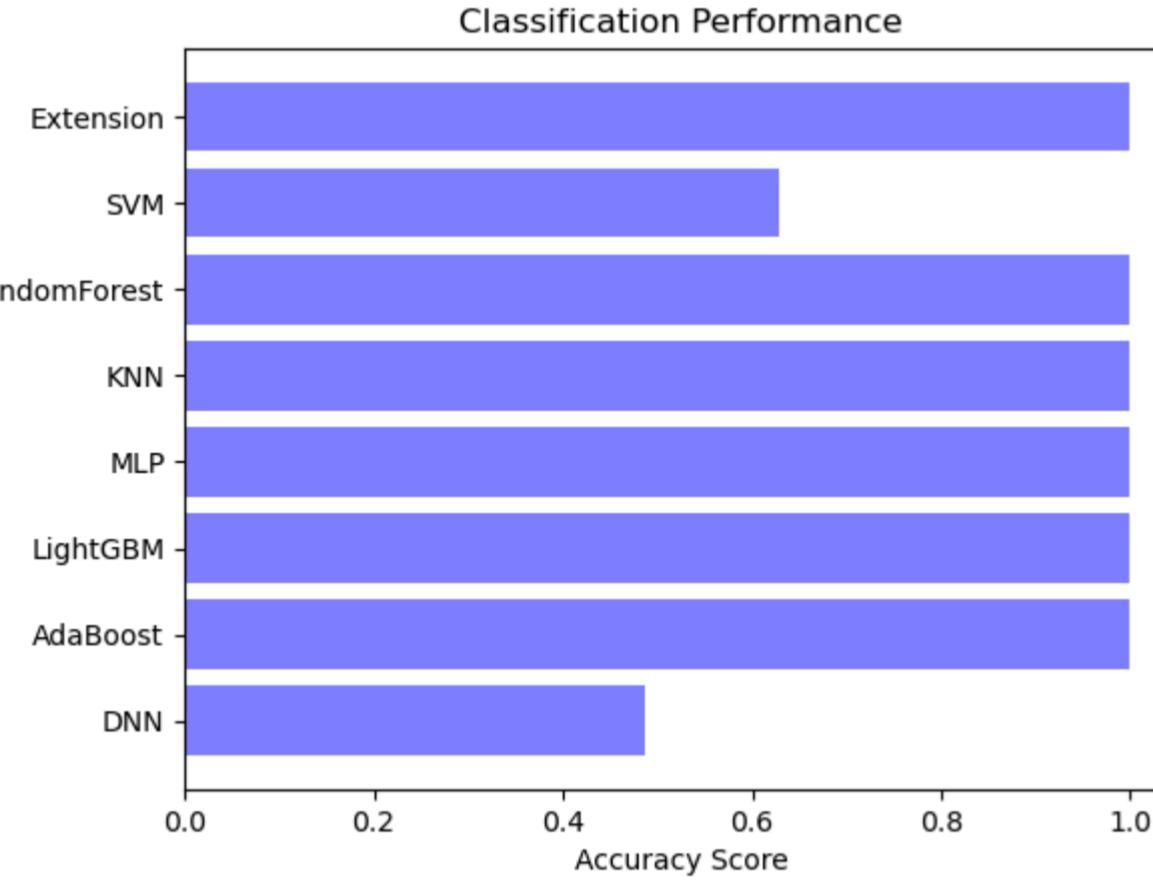
Out[34]: ['models/model\_simarg.sav']

## Graph

```
In [36]: classifier = ML_Model
y_pos = np.arange(len(classifier))
```

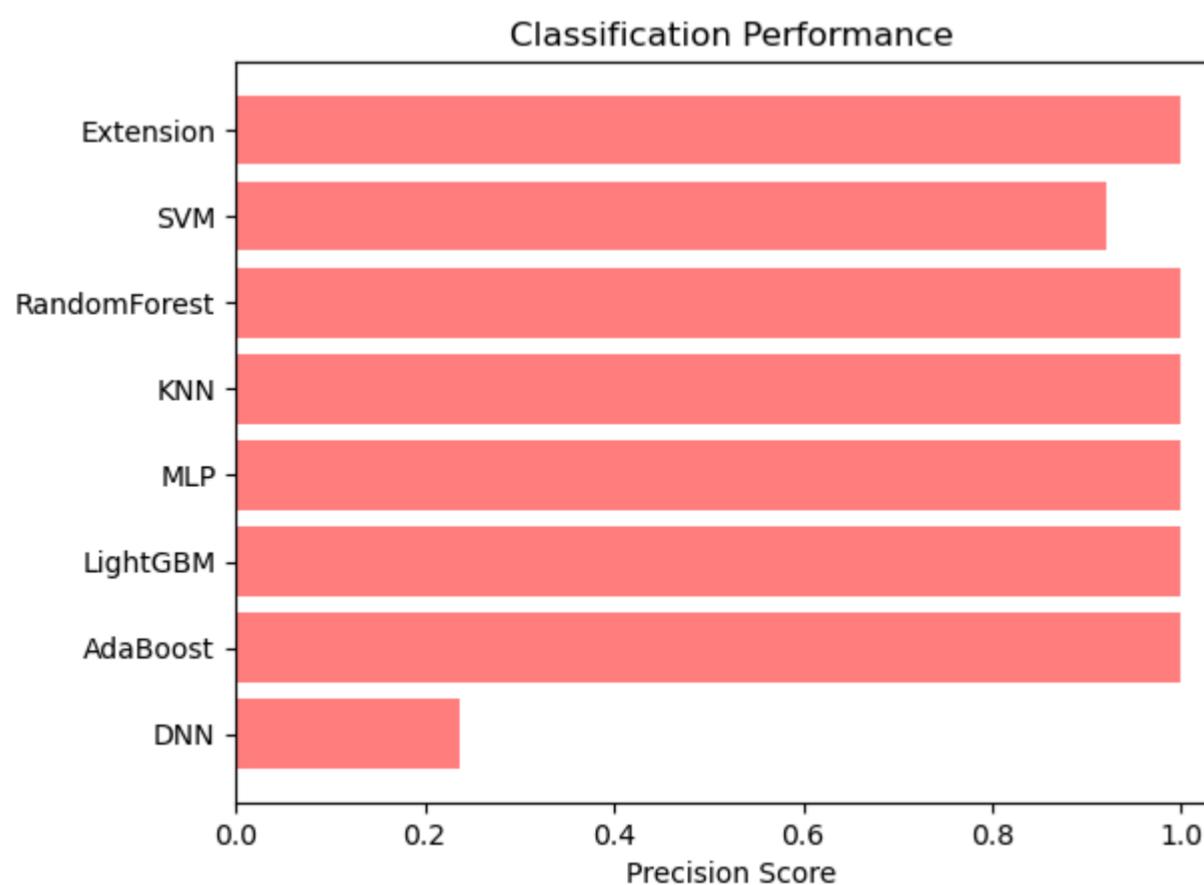
## Accuracy

```
In [37]: import matplotlib.pyplot as plt2
plt2.barh(y_pos, accuracy, align='center', alpha=0.5,color='blue')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Accuracy Score')
plt2.title('Classification Performance')
plt2.show()
```



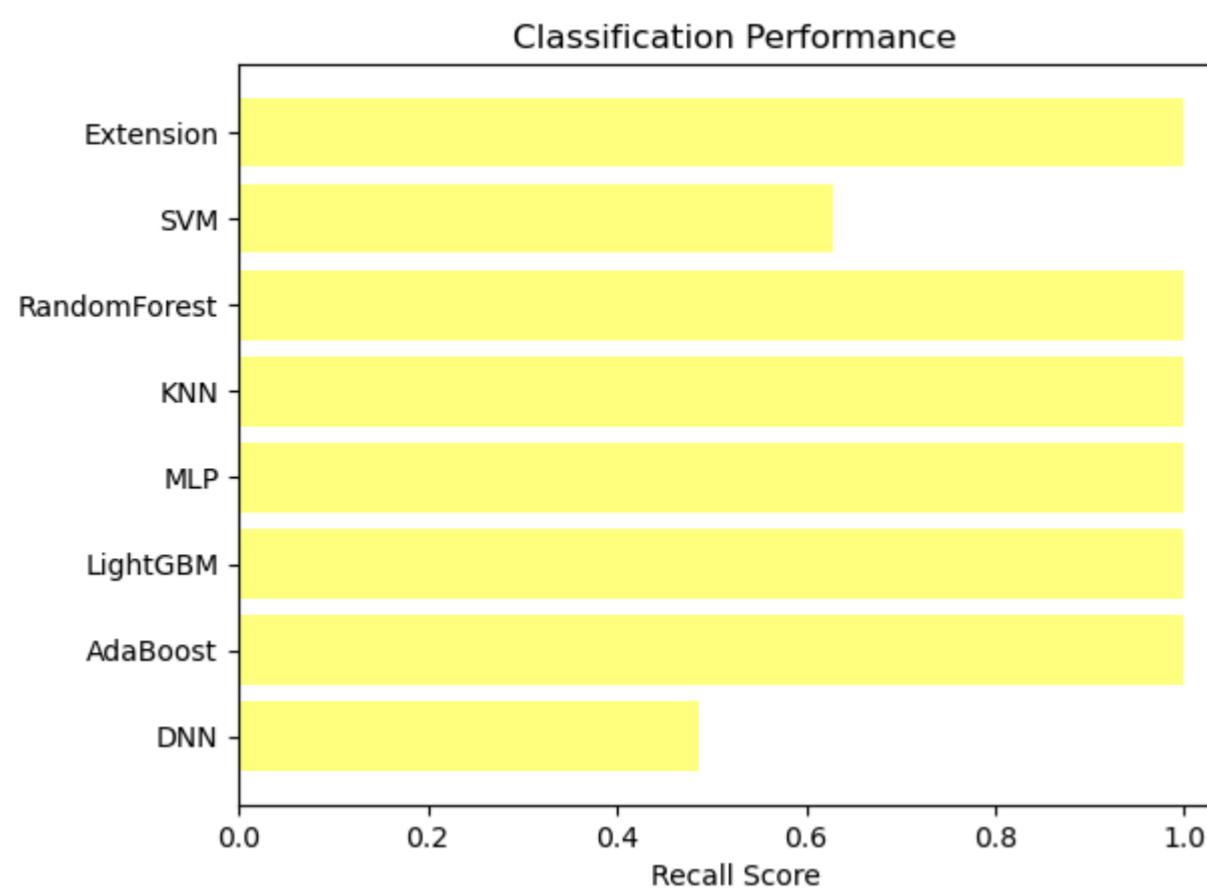
## Precision

```
In [38]: plt2.barh(y_pos, precision, align='center', alpha=0.5,color='red')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Precision Score')
plt2.title('Classification Performance')
plt2.show()
```



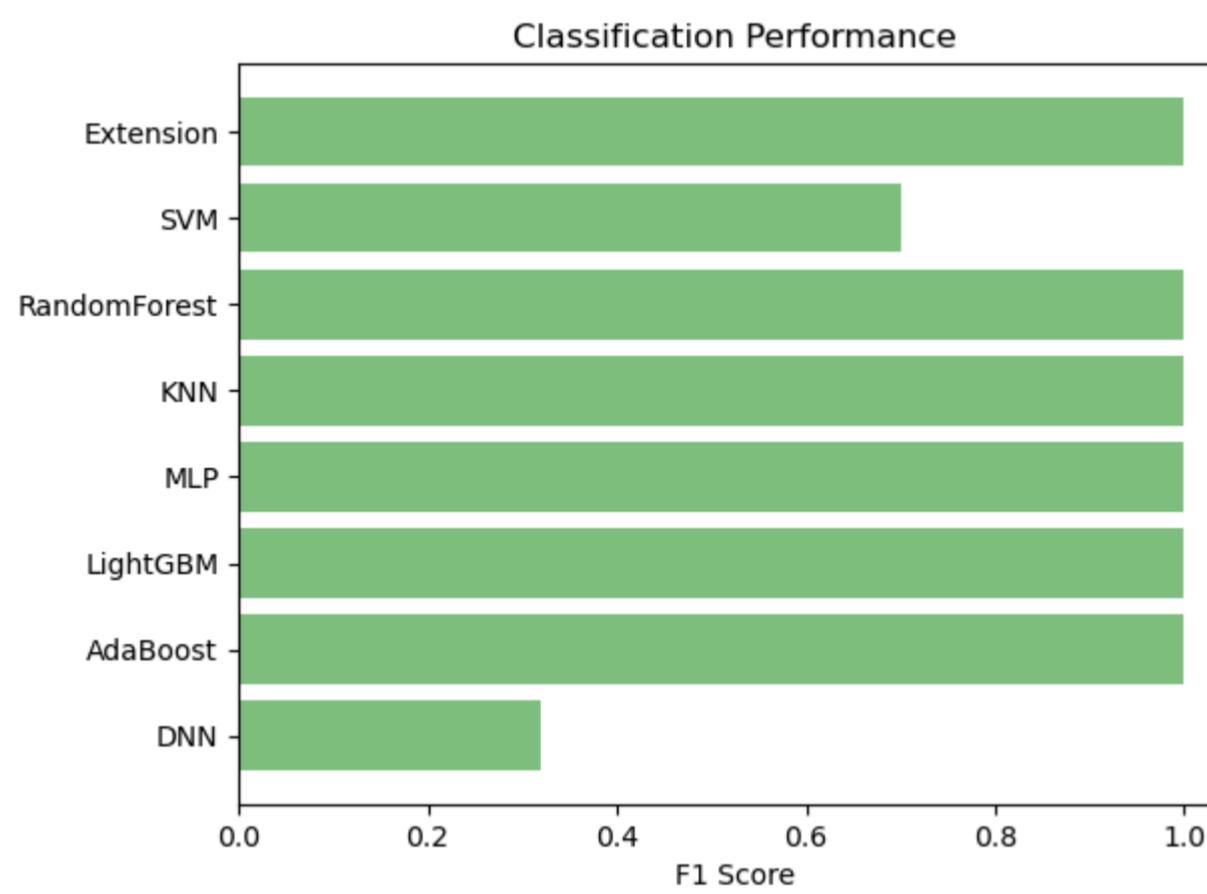
## Recall

```
In [39]: plt2.barh(y_pos, recall, align='center', alpha=0.5,color='yellow')
plt2.yticks(y_pos, classifier)
plt2.xlabel('Recall Score')
plt2.title('Classification Performance')
plt2.show()
```



## F1 Score

```
In [40]: plt2.barh(y_pos, f1score, align='center', alpha=0.5,color='green')
plt2.yticks(y_pos, classifier)
plt2.xlabel('F1 Score')
plt2.title('Classification Performance')
plt2.show()
```



```
In [ ]:
```